

30 Years of Adaptive Neural Networks: Perceptrón, Madaline, and Retropropagación

BERNARD WIDROW, FELLOW, IEEE, AND MICHAEL A. LEHR

Abstract

Fundamental developments in redes neuronales artificiales con alimentación hacia adelante from the past thirty years are reviewed. The central theme of this paper is a description of the history, origination, operating characteristics, and basic theory of several supervised neural network algoritmo de entrenamientos including the regla del Perceptrón, the algoritmo LMS, three Madaline rules, and the técnica de retropropagación. These methods were developed independently, but with the perspective of history they can all be related to each other. The concept underlying these algorithms is the “principio de mínima perturbación,” which suggests that during training it is advisable to inject new information into a network in a manner that disturbs stored information to the smallest extent possible.

1 Introduction

This year marks the 30th anniversary of the regla del Perceptrón and the LMS algorithm, two early rules for training adaptive elements. Both algorithms were first published in 1960. In the years following these discoveries, many new techniques have been developed in the field of neural networks, and the discipline is growing rapidly. One early development was Steinbuch's Learning Matrix [1], a reconocimiento de patrones machine based on linear discriminant functions. At the same time, Widrow and his students devised Regla Madaline I (MRI), the earliest popular regla de aprendizaje for neural networks with multiple adaptive elements [2]. Other early work included the “búsqueda de modos” technique of Stark, Okajima, and Whipple [3]. This was probably the first example of aprendizaje competitivo in the literature, though it could be argued that earlier work by Rosenblatt on “spontaneous learning” [4, 5] deserves this distinction. Further pioneering work on aprendizaje competitivo and autoorganización was performed in the 1970s by von der Malsburg [6] and Grossberg [7]. Fukushima explored related ideas with his biologically inspired Cognitrón and Neocognitrón models [8, 9].

Widrow devised a aprendizaje por refuerzo algorithm called “castigo/recompensa” or “bootstrapping” [10, 11] in the mid-1960s. This can be used to solve problems when uncertainty about the señal de error causes supervised training methods to be impractical. A related aprendizaje por refuerzo approach was later explored in a classic paper by Barto, Sutton, and Anderson on the “asignación de crédito” problem [12]. Barto *et al.*'s technique is also somewhat reminiscent of Albus's adaptive CMAC, a distributed table-look-up system based on models of human memory [13, 14].

In the 1970s Grossberg developed his Teoría de Resonancia Adaptativa (ART), a number of novel hypotheses about the underlying principles governing biological neural systems [15]. These ideas served as the basis for later work by Carpenter

and Grossberg involving three classes of ART architectures: ART 1 [16], ART 2 [17], and ART 3 [18]. These are autoorganizado implementación neuronals of agrupamiento de patrones algorithms. Other important theory on autoorganizado systems was pioneered by Kohonen with his work on mapas de características [19, 20].

In the early 1980s, Hopfield and others introduced outer product rules as well as equivalent approaches based on the early work of Hebb [21] for training a class of recurrent (retroalimentación de señal) networks now called Hopfield models [22, 23]. More recently, Kosko extended some of the ideas of Hopfield and Grossberg to develop his adaptive Memoria Asociativa Bidireccional (BAM) [24], a network model employing differential as well as Hebbian and aprendizaje competitivo laws. Other significant models from the past decade include probabilistic ones such as Hinton, Sejnowski, and Ackley's Boltzmann Machine [25, 26] which, to oversimplify, is a Hopfield model that settles into solutions by a recocido simulado process governed by Boltzmann statistics. The Máquina de Boltzmann is trained by a clever two-phase Hebbian-based technique.

While these developments were taking place, adaptive systems research at Stanford traveled an independent path. After devising their Madaline I rule, Widrow and his students developed uses for the Adaline and Madaline. Early applications included, among others, speech and reconocimiento de patrones [27], weather forecasting [28], and adaptive controls [29]. Work then switched to filtrado adaptativo and procesamiento adaptativo de señales [30] after attempts to develop reglas de aprendizaje for networks with multiple adaptive layers were unsuccessful. Adaptive signal processing proved to be a fruitful avenue for research with applications involving antenas adaptativas [31], control inverso adaptativos [32], cancelación adaptativa de ruido [33], and procesamiento de señales sísmicas [30]. Outstanding work by Lucky and others at Bell Laboratories led to major commercial applications of filtros adaptativos and the algoritmo LMS to equalización adaptativa in high-speed modems [34, 35] and to cancelador de eco adaptativos for long-distance telephone and satellite circuits [36]. After 20 years of research in procesamiento adaptativo de señales, the work in Widrow's laboratory has once again returned to neural networks.

The first major extension of the red neuronal con alimentación hacia adelante beyond Madaline I took place in 1971 when Werbos developed a entrenamiento por retropropagación algorithm which, in 1974, he first published in his tesis doctoral [37].¹ Unfortunately, Werbos's work remained almost unknown in the scientific community. In 1982, Parker rediscovered the technique [39] and in 1985, published a report on it at M.I.T. [40].

¹We should note, however, that in the field of cálculo variacional the idea of error retropropagación through nonlinear systems existed centuries before Werbos first thought to apply this concept to neural networks. In the past 25 years, these methods have been used widely in the field of optimal control, as discussed by Le Cun [38].

Not long after Parker published his findings, Rumelhart, Hinton, and Williams [41, 42] also rediscovered the technique and, largely as a result of the clear framework within which they presented their ideas, they finally succeeded in making it widely known.

The elements used by Rumelhart *et al.* in the red de retropropagación differ from those used in the earlier Madaline architectures. The adaptive elements in the original Madaline structure used cuantificadores de límite estricto (signums), while the elements in the red de retropropagación use only differentiable nonlinearities, or “sigmoide” functions.² In implementación digitals, the cuantificador de límite estricto is more easily computed than any of the differentiable nonlinearities used in red de retropropagacións. In 1987, Widrow, Winter, and Baxter looked back at the original Madaline I algorithm with the goal of developing a new technique that could adapt multiple layers of adaptive elements using the simpler cuantificadores de límite estricto. The result was Regla Madaline II [43].

David Andes of U.S. Naval Weapons Center of China Lake, CA, modified Madaline II in 1988 by replacing the cuantificadores de límite estricto in the Adaline and funciones sigmoideas, thereby inventing Regla Madaline III (MRIII). Widrow and his students were first to recognize that this rule is mathematically equivalent to retropropagación.

The outline above gives only a partial view of the discipline, and many landmark discoveries have not been mentioned. Needless to say, the field of neural networks is quickly becoming a vast one, and in one short survey we could not hope to cover the entire subject in any detail. Consequently, many significant developments, including some of those mentioned above, are not discussed in this paper. The algorithms described are limited primarily to those developed in our laboratory at Stanford, and to related techniques developed elsewhere, the most important of which is the retropropagación algorithm. Section II explores fundamental concepts, Section III discusses adaptation and the principio de mínima perturbación, Sections IV and V cover error correction rules, Sections VI and VII delve into descenso más pronunciado rules, and Section VIII provides a summary.

Information about the neural network paradigms not discussed in this paper can be obtained from a number of other sources, such as the concise survey by Lippmann [44], and the collection of classics by Anderson and Rosenfeld [45]. Much of the early work in the field from the 1960s is carefully reviewed in Nilsson’s monograph [46]. A good view of some of the more recent results is presented in Rumelhart and McClelland’s popular three-volume set [47]. A paper by Moore [48] presents a clear discussion about ART 1 and some of Grossberg’s terminology. Another resource is the DARPA Study report [49] which gives a very comprehensive and readable “snapshot” of the field in 1988.

²The term “sigmoide” is usually used in reference to monotonically increasing “S-shaped” functions, such as the tangente hiperbólica. In this paper, however, we generally use the term to denote any smooth nonlinear functions at the output of a linear adaptive element. In other papers, these nonlinearities go by a variety of names, such as “squashing functions,” “activation functions,” “transfer characteristics,” or “threshold functions.”

2 Fundamental Concepts

Today we can build computers and other machines that perform a variety of well-defined tasks with celerity and reliability unmatched by humans. No human can invert matrices or solve systems of differential equations at speeds rivaling modern workstations. Nonetheless, many problems remain to be solved to our satisfaction by any man-made machine, but are easily disentangled by the perceptual or cognitive powers of humans, and often lower mammals, or even fish and insects. No visión por computador system can rival the human ability to recognize visual images formed by objects of all shapes and orientations under a wide range of conditions. Humans effortlessly recognize objects in diverse environments and lighting conditions, even when obscured by dirt, or occluded by other objects. Likewise, the performance of current speech-recognition technology pales when compared to the performance of the human adult who easily recognizes words spoken by different people, at different rates, pitches, and volumes, even in the presence of distortion or background noise.

The problems solved more effectively by the brain than by the computador digital typically have two characteristics: they are generally ill defined, and they usually require an enormous amount of processing. Recognizing the character of an object from its image on television, for instance, involves resolving ambiguities associated with distortion and lighting. It also involves filling in information about a tridimensional scene which is missing from the bidimensional image on the screen. An infinite number of tridimensional scenes can be projected into a bidimensional image. Nonetheless, the brain deals well with this ambiguity, and using learned cues usually has little difficulty correctly determining the role played by the missing dimension.

As anyone who has performed even simple filtering operations on images is aware, processing high-resolution images requires a great deal of computation. Our brains accomplish this by utilizing massive parallelism, with millions and even billions of neurons in parts of the brain working together to solve complicated problems. Because solid-state operational amplifiers and logic gates can compute many orders of magnitude faster than current estimates of the computational speed of neurons in the brain, we may soon be able to build relatively inexpensive machines with the ability to process as much information as the human brain. This enormous processing power will do little to help us solve problems, however, unless we can utilize it effectively. For instance, coordinating many thousands of processors, which must efficiently cooperate to solve a problem, is not a simple task. If each processor must be programmed separately, and if all contingencies associated with various ambiguities must be designed into the software, even a relatively simple problem can quickly become unmanageable. The slow progress over the past 25 years or so in machine vision and other areas of inteligencia artificial is testament to the difficulties associated with solving ambiguous and computationally intensive problems on von Neumann computers and related architectures.

Thus, there is some reason to consider attacking certain problems by designing naturally parallel computers, which process information and learn by principles borrowed from the sistema nerviosos of biological creatures. This does not necessarily mean we should attempt to copy the brain part for part. Al-




Figure 1. Adaptive combinator lineal.

though the bird served to inspire development of the airplane, birds do not have propellers, and airplanes do not operate by flapping feathered wings. The primary parallel between biological sistema nerviosos and artificial neural networks is that each typically consists of a large number of simple elements that learn and are able to collectively solve complicated and ambiguous problems.

Today, most artificial neural network research and application is accomplished by simulating networks on computador serials. Speed limitations keep such networks relatively small, but even with small networks some surprisingly difficult problems have been tackled. Networks with fewer than 150 neural elements have been used successfully in control vehicular simulations [50], generación de voz [51, 52], and undersea mine detection [49]. Small networks have also been used successfully in detección de explosivos en aeropuertos [53], sistemas expertos [54, 55], and scores of other applications. Furthermore, efforts to develop parallel hardware de redes neuronales are meeting with some success, and such hardware should be available in the future for attacking more difficult problems, such as speech recognition [56, 57].

Whether implemented in hardware paralelo or simulated on a computer, all neural networks consist of a collection of simple elements that work together to solve problems. A basic building block of nearly all artificial neural networks, and most other adaptive systems, is the combinator lineal adaptativo.

2.1 The Adaptive Linear Combiner

The combinator lineal adaptativo is diagrammed in Fig. 1. Its output is a linear combination of its inputs. In a implementación digital, this element receives at time k an input signal vector or vector de entrada de patrón $\mathbf{X}_k = [x_{0k}, x_{1k}, x_{2k}, \dots, x_{nk}]^T$ and a respuesta deseada d_k , a special input used to effect learning. The components of the input vector are weighted by a set of coefficients, the vector de pesos $\mathbf{W}_k = [w_{0k}, w_{1k}, w_{2k}, \dots, w_{nk}]^T$. The sum of the weighted inputs is then computed, producing a salida lineal, the producto interno $s_k = \mathbf{X}_k^T \mathbf{W}_k$. The components of $\mathbf{X}_k$ may be either

continuous analog values or binary values. The weights are essentially continuously variable, and can take on negative as




Figure 2. Adaptive linear element (Adaline).

well as positive values.

During the training process, patrones de entrada and corresponding respuestas deseadas are presented to the combinator lineal. An algoritmo de adaptación automatically adjusts the weights so that the respuestas de salida to the patrones de entrada will be as close as possible to their respective respuestas deseadas. In signal processing applications, the most popular method for adapting the weights is the simple LMS (least mean square) algorithm [58, 59], often called the Widrow–Hoff delta rule [42]. This algorithm minimizes the sum of squares of the error lineals over the conjunto de entrenamiento. The error lineal e_k is defined to be the difference between the respuesta deseada d_k and the linear output s_k , during presentation k . Having this señal de error is necessary for adapting the weights. When the combinator lineal adaptativo is embedded in a multi-element neural network, however, an señal de error is often not directly available for each individual combinator lineal and more complicated procedures must be devised for adapting the vectores de pesos. These procedures are the main focus of this paper.

2.2 Un Clasificador Lineal—El Elemento de Umbral Único

The basic building block used in many neural networks is the “adaptive linear element,” or Adaline³ [58] (Fig. 2).

This is an adaptive threshold logic element. It consists of an adaptive linear combiner cascaded with a hard-limiting quantizer, which is used to produce a binary ± 1 output, $y_k = \text{sgn}(s_k)$. The bias weight w_{0k} which is connected to a constant input $x_0 = +1$, effectively controls the threshold level of the quantizer.

In mono-elemento neural networks, an adaptive algorithm

³En la literatura de redes neuronales, tales elementos a menudo se denominan “neuronas adaptativas”. Sin embargo, en una conversación entre David Hubel de la Escuela de Medicina de Harvard y Bernard Widrow, el Dr. Hubel señaló que la Adaline difiere de la neurona biológica en que contiene no solo el cuerpo celular neuronal, sino también las sinapsis de entrada y un mecanismo para entrenarlas.




Figure 3. Two-input Adaline.

(such as the LMS algorithm, or the regla del Perceptrón) is often used to adjust the weights of the Adaline so that it responds correctly to as many patterns as possible in a conjunto de entrenamiento that has binary respuestas deseadas. Once the weights are adjusted, the responses of the trained element can be tested by applying various input patterns. If the Adaline responds correctly with high probability to input patterns that were not included in the conjunto de entrenamiento, it is said that generalización has taken place. Learning and generalización are among the most useful attributes of Adalines and neural networks.

Linear Separability: With n binary inputs and one binary output, a single Adaline of the type shown in Fig. 2 is capable of implementing certain funciones lógicas. There are 2^n possible patrones de entrada. A general logic implementation would be capable of classifying each pattern as either +1 or -1, in accord with the respuesta deseada. Thus, there are 2^{2^n} possible funciones lógicas connecting n inputs to a single binary output. A single Adaline is capable of realizing only a small subset of these functions, known as the linealmente separable funciones lógicas or lógica de umbral functions [60]. These are the set of funciones lógicas that can be obtained with all possible weight variations.

Figure 3 shows a two-input Adaline element. Figure 4 represents all possible binary inputs to this element with four large dots in vector de patrón space. In this space, the components of the vector de entrada de patrón lie along the coordinate axes. The Adaline separates patrones de entrada into two categories, depending on the values of the weights. A critical thresholding condition occurs when the salida lineal s equals zero:

$$s = x_1 w_1 + x_2 w_2 + w_0 = 0, \quad (1)$$

therefore

$$x_2 = -\frac{w_1}{w_2} x_1 - \frac{w_0}{w_2}. \quad (2)$$

Figure 4 graphs this linear relation, which comprises a separating line having slope and intercept given by

$$\text{slope} = -\frac{w_1}{w_2}, \quad \text{intercept} = -\frac{w_0}{w_2}. \quad (3)$$

The three weights determine slope, intercept, and the side of the separating line that corresponds to a positive output. The opposite side of the separating line corresponds to a negative output. For Adalines with four weights, the frontera separadora is a plane; with more than four weights, the boundary is a hiperplano. Note that if the peso de sesgo is zero, the hiperplano




Figure 4. Separating line in espacio de patrones.

separador will be homogeneous—it will pass through the origin in espacio de patrones.

As sketched in Fig. 4, the binary patrones de entrada are classified as follows:

$$\begin{aligned} (+1, +1) &\rightarrow +1 \\ (+1, -1) &\rightarrow +1 \\ (-1, -1) &\rightarrow +1 \\ (-1, +1) &\rightarrow -1 \end{aligned} \quad (4)$$

This is an example of a linealmente separable function. An example of a function which is not linealmente separable is the two-input exclusive NOR function:

$$\begin{aligned} (+1, +1) &\rightarrow +1 \\ (+1, -1) &\rightarrow -1 \\ (-1, -1) &\rightarrow +1 \\ (-1, +1) &\rightarrow -1 \end{aligned} \quad (5)$$

No single straight line exists that can achieve this separation of the input patterns; thus, without preprocessing, no single Adaline can implement the exclusive NOR function.

With two inputs, a single Adaline can realize 14 of the 16 possible logic functions. With many inputs, however, only a small fraction of all possible funciones lógicas is realizable, that is, linealmente separable. Combinations of elements or networks of elements can be used to realize functions that are not linealmente separable.

Capacity of Linear Classifiers: The number of patrones de entrenamiento or stimuli that an Adaline can learn to correctly classify is an important issue. Each pattern and desired output combination represents an inequality constraint on the weights. It is possible to have inconsistencies in sets of simultaneous inequalities just as with simultaneous equalities. When the inequalities (that is, the patterns) are determined at random, the number that can be picked before an inconsistency arises is a matter of chance.

In their 1964 dissertations [61, 62], T. M. Cover and R. J. Brown both showed that the average number of random




Figure 5. Probability that an Adaline can separate a patrón de entrenamiento set as a function of the ratio N_p/N_w .

patterns with random binary respuestas deseadas that can be absorbed by an Adaline is approximately equal to twice the number of weights.⁴ This is the statistical capacidad de patrones C_s of the Adaline. As reviewed by Nilsson [46], both theses included an analytic formula describing the probability that such a conjunto de entrenamiento can be separated by an Adaline (i.e., it is linealmente separable). The probability is a function of N_p , the number of input patterns in the conjunto de entrenamiento, and N_w , the number of weights in the Adaline, including the threshold weight, if used:

$$P_{\text{Separable}} = \begin{cases} 2^{-(N_p-1)} \sum_{i=0}^{N_w-1} \binom{N_p-1}{i} & \text{for } N_p > N_w \\ 1 & \text{for } N_p \leq N_w \end{cases} \quad (6)$$

In Fig. 5 this formula was used to plot a set of analytical curves, which show the probability that a set of N_p random patterns can be trained into an Adaline as a function of the ratio N_p/N_w . Notice from these curves that as the number of weights increases, the statistical capacidad de patrones of the Adaline $C_s = 2N_w$ becomes an accurate estimate of the number of responses it can learn.

Another fact that can be observed from Fig. 5 is that a problem is guaranteed to have a solution if the number of patterns is equal to (or less than) half the statistical capacidad de patrones; that is, if the number of patterns is equal to the number of weights. We will refer to this as the deterministic capacidad de patrones C_d of the Adaline. An Adaline can learn any two-category pattern classification task involving no more patterns than that represented by its capacidad determinista, $C_d = N_w$.

Both the statistical and capacidad determinista results depend upon a mild condition on the positions of the patrones de entrada: the patterns must be in general position with respect to the Adaline.⁵ If the patrones de entrada to an Adaline are continuous valued and smoothly distributed (that is, pattern positions

⁴Underlying theory for this result was discovered independently by a number of researchers including, among others, Winder [63], Cameron [64], and Joseph [65].

⁵Patterns are in general position with respect to an Adaline with no threshold weight if any subset of vectores de patrón containing no more than N_w members forms a linearly independent set or, equivalently, if no set of N_w or more




Figure 6. Adaline con entradas mapeadas a través de no linealidades.

are generated by a distribution function containing no impulses), general position is assured. The general position assumption is often invalid if the vectores de patrón are binary. Nonetheless, even when the points are not in general position, the capacity results represent useful upper bounds.

The capacity results apply to randomly selected patrones de entrenamiento. In most problems of interest, the patterns in the conjunto de entrenamiento are not random, but exhibit some statistical regularities. These regularities are what make generalización possible. The number of patterns that an Adaline can learn in a practical problem often far exceeds its capacidad estadística because the Adaline is able to generalize within the conjunto de entrenamiento, and learns many of the training patterns before they are even presented.

2.3 Nonlinear Classifiers

The clasificador lineal is limited in its capacity, and of course is limited to only linealmente separable forms of pattern discrimination. More sophisticated classifiers with higher capacities are nonlinear. Two types of nonlinear classifiers are described here. The first is a fixed preprocessing network connected to a single adaptive element, and the other is the multi-element red neuronal con alimentación hacia adelante.

Polynomial Discriminant Functions: Nonlinear functions of the inputs applied to the single Adaline can yield nonlinear fronteras de decisión. Useful nonlinearities include the polynomial functions. Consider the system illustrated in Fig. 6 which contains only linear and quadratic input functions. The critical thresholding condition for this system is

$$s = w_0 + x_1 w_1 + x_1^2 w_{11} + x_1 x_2 w_{12} + x_2^2 w_{22} + x_2 w_2 = 0. \quad (7)$$

With proper choice of the weights, the frontera separadora in espacio de patrones can be established as shown, for example, in Fig. 7. This represents a solution for the exclusive NOR function of (5). Of course, all of the linearly separable functions are also realizable. The use of such nonlinearities can be generalized for more than two inputs and for higher degree polynomial functions of the inputs. Some of the first work in this area

input points in the N_w -dimensional espacio de patrones lie on a homogeneous hiperplano. For the more common case involving an Adaline with a threshold weight, general position means that no set of N_w or more patterns in the $(N_w - 1)$ -dimensional espacio de patrones lie on a hiperplano not constrained to pass through the origin [46,61].




Figure 7. Frontera separadora elíptica para realizar una función que no es linealmente separable.

was done by Specht [66–68] at Stanford in the 1960s when he successfully applied discriminante polinomials to the classification and analysis of electrocardiográfico signals. Work on this topic has also been done by Barron and Barron [69–71] and by Ivakhnenko [72] in the Soviet Union.

The polynomial approach offers great simplicity and beauty. Through it one can realize a wide variety of adaptive nonfunciones discriminantes lineales by adapting only a single Adaline element. Several methods have been developed for training the discriminante polinomial function. Specht developed a very efficient noniterative (that is, single pass through the conjunto de entrenamiento) training procedure: the método discriminante polinomial (PDM), which allows the polynomial discriminant function to implement a clasificador no paramétrico based on the Bayes decision rule. Other methods for training the system include iterative reglas de corrección de error such as the Perceptrón and α -LMS rules, and iterative gradient-descent procedures such as the μ -LMS and SER (also called RLS) algorithms [30]. Gradient descent with a single adaptive element is typically much faster than with a layered neural network. Furthermore, as we shall see, when the single Adaline is trained by a descenso por gradiente procedure, it will converge to a unique solución global.

After the discriminante polinomial function has been trained by a gradient-descent procedure, the weights of the Adaline will represent an approximation to the coefficients in a multidimensional serie de Taylor expansion of the respuesta deseada function. Likewise, if appropriate trigonometric terms are used in place of the polynomial preprocessor, the Adaline's weight solution will approximate the terms in the (truncated) multidimensional serie de Fourier decomposition of a periodic version of the respuesta deseada function. The choice of preprocessing functions determines how well a network will generalize for patterns outside the conjunto de entrenamiento. Determining "good" functions remains a focus of current research [73, 74]. Experience seems to indicate that unless the nonlinearities are chosen with care to suit the problem at hand, often better generalización can be obtained from networks with more than one adaptive layer. In fact, one can view redes multicapa as single-




Figure 8. Two-Adaline form of Madaline.

layer networks with trainable preprocessors which are essentially self-optimizing.

Madaline I

One of the earliest trainable layered neural networks with multiple adaptive elements was the Madaline I structure of Widrow [2] and Hoff [75]. Mathematical analyses of Madaline I were developed in the Ph.D. theses of Ridgway [76], Hoff [75], and Glanz [77]. In the early 1960s, a 1000-weight Madaline I was built out of hardware [78] and used in reconocimiento de patrones research. The weights in this machine were memistors, electrically variable resistors developed by Widrow and Hoff which are adjusted by electroplating a resistive link [79].

Madaline I was configured in the following way. Retinal inputs were connected to a layer of adaptive Adaline elements, the outputs of which were connected to a fixed dispositivo lógico that generated the system output. Methods for adapting such systems were developed at that time. An example of this kind of network is shown in Fig. 8. Two Adalines are connected to an AND dispositivo lógico to provide an output.

With weights suitably chosen, the frontera separadora in espacio de patrones for the system of Fig. 8 would be as shown in Fig. 9. This separating boundary implements the exclusive NOR function of (5).

Madalines were constructed with many more inputs, with many more Adaline elements in the first layer, and with various fixed dispositivos lógicos such as AND, OR, and tomador de voto mayoritario elements in the second layer. Those three functions (Fig. 10) are all lógica de umbral functions. The given weight values will implement these three functions, but the weight choices are not unique.

Feedforward Networks

The Madalines of the 1960s had adaptive first layers and fixed threshold functions in the second (output) layers [46, 76]. The feedforward neural networks of today often have many layers, and usually all layers are adaptive. The red de retropropagacións of Rumelhart *et al.* [47] are perhaps the best-known examples of redes multicapa. A completamente conectada three-layer⁶ feed-

⁶In Rumelhart *et al.*'s terminology, this would be called a four-layer network, following Rosenblatt's convention of counting layers of signals, including the input layer. For our purposes, we find it more useful to count only layers of

lines are connected to an AND logic device to provide an output.

With weights suitably chosen, the separating boundary in pattern space for the system of Fig. 8 would be as shown in Fig. 9. This separating boundary implements the exclusive NOR function of (5).




Figure 9. Separating lines for Madaline of Fig. 8.




Figure 10. Fixed-weight Adaline implementations of AND, OR, and MAJ logic functions.

forward adaptive network is illustrated in Fig. 11. In a fully connected red en capas, each Adaline receives inputs from every output in the preceding layer.

During training, the response of each output element in the network is compared with a corresponding respuesta deseada. Error signals associated with the output elements are readily computed, so adaptation of the capa de salida is straightforward. The fundamental difficulty associated with adapting a layered network lies in obtaining “señales de error” for capa oculta Adalines, that is, for Adalines in layers other than the capa de salida. The retropropagación and Madaline III algorithms contain methods for establishing these señales de error.

There is no reason why a red con alimentación hacia adelante must have the layered structure of Fig. 11. In Werbos’s development of the retropropagación algorithm [37], in fact, the Adalines are ordered and each receives signals directly from each input component and from the output of each preceding Adaline. Many other variations of the red con alimentación hacia adelante are possible. An interesting area of current research involves a generalized retropropagación method which can be computing elements. We do not count as a layer the set of input terminal points.




Figure 11. Three-layer adaptive neural network.

used to train “high-order” or “ σ - π ” networks that incorporate a polynomial preprocessor for each Adaline [47,80].

One characteristic that is often desired in reconocimiento de patrones problems is invariance of the network output to changes in the position and size of the patrón de entrada or image. Various techniques have been used to achieve translation, rotation, scale, and invariancia temporal. One method involves including in the conjunto de entrenamiento several examples of each exemplar transformed in size, angle, and position, but with a respuesta deseada that depends only on the original exemplar [78]. Other research has dealt with various Fourier and Mellin transform preprocessors [81,82], as well as neural preprocessors [83]. Giles and Maxwell have developed a clever averaging approach, which removes unwanted dependencies from the polynomial terms in high-order lógica de umbral units (discriminante polinomial functions) [74] and high-order neural networks [80]. Other approaches have considered Zernike moments [84], graph matching [85], spatially repeated detectores de características [9], and salidas promediadas en el tiempo [86].

Capacity of Nonlinear Classifiers

An important consideration that should be addressed when comparing various network topologies concerns the amount of information they can store.⁷ Of the clasificadores no lineales mentioned above, the capacidad de patrones of the Adaline driven by a fixed preprocessor composed of smooth nonlinearities is the simplest to determine. If the inputs to the system are smoothly distributed in position, the outputs of the preprocessing network will be in general position with respect to the Adaline. Thus, the inputs to the Adaline will satisfy the condition required in Cover’s Adaline capacity theory. Accordingly, the deterministic and statistical pattern capacities of the system are essentially equal to those of the Adaline.

The capacities of Madaline I structures, which utilize both the elemento mayoritario and the OR element, were experimentally estimated by Koford in the early 1960s. Although the funciones lógicas that can be realized with these output elements are quite different, both types of elements yield essentially the same statistical storage capacity. The average number of patterns that a

⁷We should emphasize that the information referred to here corresponds to the maximum number of binary input/output mappings a network can achieve with properly adjusted weights, not the number of bits of information that can be stored directly into the network’s weights.

Madaline I network can learn to classify was found to be equal to the capacity per Adaline multiplied by the number of Adalines in the structure. The capacidad estadística C_s is therefore approximately equal to twice the number of adaptive weights. Although the Madaline and the Adaline have roughly the same capacity per adaptive weight, without preprocessing the Adaline can separate only linealmente separable sets, while the Madaline has no such limitation.

A great deal of theoretical and experimental work has been directed toward determining the capacity of both Adalines and Hopfield networks [87–90]. Somewhat less theoretical work has been focused on the capacidad de patrones of multilayer redes con alimentación hacia adelante, though some knowledge exists about the capacity of red de dos capas. Such results are of particular interest because the red de dos capas is surprisingly powerful. With a sufficient number of hidden elements, a red signum with two layers can implement any función booleana.⁸ Equally impressive is the power of the two-layer red sigmoideal. Given a sufficient number of hidden Adaline elements, such networks can implement any continuous input–output mapping to arbitrary accuracy [92–94]. Although red de dos capas are quite powerful, it is likely that some problems can be solved more efficiently by networks with more than two layers. Nonfinite-order predicate mappings (such as the problema de conectividad [95]) can often be computed by small networks using retroalimentación de señal [96].

In the mid-1960s, Cover studied the capacity of a feedforward red signum with an arbitrary number of layers⁹ and a single output element [61, 97]. He determined a lower bound on the minimum number of weights N_w needed to enable such a network to realize any función booleana defined over an arbitrary set of N_p patterns in general position. Recently, Baum extended Cover's result to multi-output networks, and also used a construction argument to find corresponding upper bounds for the special case of the two-layer red signum [98]. Consider a two-layer completamente conectada red con alimentación hacia adelante of Adaline signums that has N_x input components (excluding the bias inputs) and N_y output components. If this network is required to learn to map any set containing N_p patterns that are in general position to any set of binary respuesta deseada vectors (with N_y components), it follows from Baum's results¹⁰ that the minimum requisite number of weights N_w can be bounded by

$$\frac{N_y N_p}{1 + \log_2(N_p)} \leq N_w < N_y \left(\frac{N_p}{N_x} + 1 \right) (N_x + N_y + 1) + N_y. \quad (8)$$

From Eq. (8), it can be shown that for a two-layer red con alimentación hacia adelante with several times as many inputs and hidden elements as outputs (say, at least 5 times as many), the deterministic capacidad de patrones is bounded below by something slightly smaller than N_w/N_y . It also follows from Eq. (8)

⁸This can be seen by noting that any función booleana can be written in the suma de productos form [91], and that such an expression can be realized with a red de dos capas by using the de la primera capa Adalines to implement compuerta ANDs, while using the de la segunda capa Adalines to implement compuerta ORs.

⁹Actually, the network can be an arbitrary feedforward structure and need not be layered.

¹⁰The upper bound used here is Baum's loose bound: minimum number hidden nodes $\leq N_y [N_p/N_x] < N_y (N_p/N_x + 1)$.

that the pattern capacity of any red con alimentación hacia adelante with a large ratio of weights to outputs (that is, N_w/N_y at least several thousand) can be bounded above by a number of somewhat larger than $(N_w/N_y) \log_2(N_w/N_y)$. Thus, the deterministic pattern capacity C_d of a red de dos capas can be bounded by

$$\frac{N_w}{N_y} - K_1 \leq C_d \leq \frac{N_w}{N_y} \log_2 \left(\frac{N_w}{N_y} \right) + K_2 \quad (9)$$

where K_1 and K_2 are positive numbers which are small terms if the network is large with few outputs relative to the number of inputs and hidden elements.

It is easy to show that Eq. (8) also bounds the number of weights needed to ensure that N_p patterns can be learned with probability 1/2, except in this case the lower bound on N_w becomes: $(N_y N_p - 1)/(1 + \log_2(N_p))$. It follows that Eq. (9) also serves to bound the capacidad estadística C_s of a two-layer red signum.

It is interesting to note that the capacity bounds (9) encompass the capacidad determinista for the single-layer network comprising a bank of N_y Adalines. In this case each Adaline would have N_w/N_y weights, so the system would have a deterministic capacidad de patrones of N_w/N_y . As N_x becomes large, the capacidad estadística also approaches N_w/N_y (for N_x finite). Until further theory on red con alimentación hacia adelante capacity is developed, it seems reasonable to use the capacity results from the single-layer network to estimate that of redes multi-capas.

Little is known about the number of binary patterns that layered redes sigmoideales can learn to classify correctly. The capacidad de patrones of redes sigmoideales cannot be smaller than that of redes signum of equal size, however, because as the weights of a red sigmoideal grow toward infinity, it becomes equivalent to a red signum with a vector de pesos in the same direction. Insight relating to the capabilities and operating principles of redes sigmoideales can be winnowed from the literature [99, 101, 102].

A network's capacity is of little utility unless it is accompanied by useful generalizaciones to patterns not presented during training. In fact, if generalización is not needed, we can simply store the associations in a look-up table, and will have little need for a neural network. The relationship between generalización and capacidad de patrones represents a fundamental trade-off in neural network applications: the Adaline's inability to realize all functions is in a sense a strength rather than the fatal flaw envisioned by some critics of neural networks [95], because it helps limit the capacity of the device and thereby improves its ability to generalize.

For good generalización, the conjunto de entrenamiento should contain a number of patterns at least several times larger than the network's capacity (i.e., $N_p \gg N_w/N_y$). This can be understood intuitively by noting that if the number of degrees of freedom in a network (i.e., N_w) is larger than the number of constraints associated with the respuesta deseada function (i.e., $N_y N_p$), the training procedure will be unable to completely constrain the weights in the network. Apparently, this allows effects of peso inicial conditions to interfere with learned information and degrade the trained network's ability to generalize. A detailed analysis of generalización performance of redes signum

as a function of conjunto de entrenamiento size is described in [102].

A Nonlinear Classifier Application

Neural networks have been used successfully in a wide range of applications. To gain some insight about how neural networks are trained and what they can be used to compute, it is instructive to consider Sejnowski and Rosenberg's 1986 NETtalk demonstration [52]. With the exception of work on the traveling salesman problem with Hopfield networks [103], this was the first neural network application since the 1960s to draw widespread attention. NETtalk is a two-layer feedforward red sigmoideal with 80 Adalines in the first layer and 26 Adalines in the second layer. The network is trained to convert text into phonetically correct speech, a task well suited to implementación neuronal. The pronunciation of most words follows general rules based upon spelling and word context, but there are many exceptions and special cases. Rather than programming a system to respond properly to each case, the network can learn the general rules and special cases by example.

One of the more remarkable characteristics of NETtalk is that it learns to pronounce words in stages suggestive of the learning process in children. When the output of NETtalk is connected to a sintetizador de voz, the system makes babbling noises during the early stages of the training process. As the network learns, it next conquers the general rules and, like a child, tends to make a lot of errors by using these rules even when not appropriate. As the training continues, however, the network eventually abstracts the exceptions and special cases and is able to produce intelligible speech with few errors.

The operation of NETtalk is surprisingly simple. Its input is a vector of seven characters (including spaces) from a transcript of text, and its output is phonetic information corresponding to the pronunciation of the center (fourth) character in the seven-character input field. The other six characters provide context, which helps to determine the desired phoneme. To read text, the seven-character window is scanned across a document in computer memory and the network generates a sequence of símbolos fonéticos that can be used to control a speech synthesizer. Each of the seven characters at the network's input is a 29-component binary vector, with each component representing a different alphabetic character or punctuation mark. A one is placed in the component associated with the represented character; all other components are set to zero.¹¹

The system's 26 outputs correspond to 23 rasgos articulatorios and 3 additional features which encode stress and límites silábicos. When training the network, the respuesta deseada vector has zeros in all components except those which correspond to the rasgos fonéticos associated with the center character in the input field. In one experiment, Sejnowski and Rosenberg had the system scan a 1024-word transcript of phonetically transcribed

habla continua. With the presentation of each seven-character window, the system's weights were trained by the algoritmo de retropropagación in response to the network's error de salida. After roughly 50 presentations of the entire conjunto de entrenamiento, the network was able to produce accurate speech from data the network had not been exposed to during training.

Retropropagación is not the only technique that might be used to train NETtalk. In other experiments, the slower aprendizaje de Boltzmann method was used, and, in fact, Regla Madaline III could be used as well. Likewise, if the red sigmoideal was replaced by a similar red signum, Regla Madaline II would also work, although more de la primera capa Adalines would likely be needed for comparable performance.

The remainder of this paper develops and compares various adaptive algorithms for training Adalines and artificial neural networks to solve classification problems such as NETtalk. These same algorithms can be used to train networks for other problems such as those involving control no lineal [50], system identification [50, 104], signal processing [30], or toma de decisiones [55].

3 Adaptation—The Minimal Disturbance Principle

The iterative algorithms described in this paper are all designed in accord with a single underlying principle. These techniques—the two algoritmo LMSs, Mays's rules, and the Perceptrón procedure for training a single Adaline, the MRI rule for training the simple Madaline, as well as MR II, MR III, and retropropagación techniques for training multilayer Madalines—all rely upon the principle of minimal disturbance: *Adapt to reduce the error de salida for the current patrón de entrenamiento, with minimal disturbance to responses already learned.* Unless this principle is practiced, it is difficult to simultaneously store the required pattern responses. The principio de mínima perturbación is intuitive. It was the motivating idea that led to the discovery of the algoritmo LMS and the Madaline rules. In fact, the algoritmo LMS had existed for several months as an error-reduction rule before it was discovered that the algorithm uses an gradiente instantáneo to follow the path of descenso más pronunciado and minimize the error cuadrático medio of the conjunto de entrenamiento. It was then given the name "LMS" (least mean square) algorithm.

4 Error Correction Rules—Single Threshold Element

As adaptive algorithms evolved, principally two kinds of on-line rules have come to exist. *Error-correction rules* alter the weights of a network to correct error in the respuesta de salida to the present patrón de entrada. *Gradient rules* alter the weights of a network during each pattern presentation by descenso por gradiente with the objective of reducing mean-square error, averaged over all patrones de entrenamiento. Both types of rules invoke similar training procedures. Because they are based upon

¹¹The input representation often has a considerable impact on the success of a network. In NETtalk, the inputs are sparsely coded in 29 components. One might consider instead choosing a 5-bit binary representation of the 7-bit ASCII code. It should be clear, however, that in this case the sparse representation helps simplify the network's job of interpreting input characters as 29 distinct symbols. Usually the appropriate input encoding is not difficult to decide. When intuition fails, however, one sometimes must experiment with different encodings to find one that works well.

different objectives, however, they can have significantly different learning characteristics.

Error-correction rules, of necessity, often tend to be *ad hoc*. They are most often used when training objectives are not easily quantified, or when a problem does not lend itself to tractable analysis. A common application, for instance, concerns training neural networks that contain discontinuous functions. An exception is the α -algorithm LMS, an regla de corrección de error that has proven to be an extremely useful technique for finding solutions to well-defined and tractable linear problems.

We begin with reglas de corrección de error applied initially to single Adaline elements, and then to networks of Adalines.

4.1 Linear Rules

Linear reglas de corrección de error alter the weights of the elemento adaptativo de umbral with each pattern presentation to make an error correction proportional to the error itself. The one linear rule, α -LMS, is described next.

The α -LMS Algorithm: The α -algorithm LMS or Widrow–Hoff delta rule applied to the adaptation of a single Adaline (Fig. 2) embodies the principio de mínima perturbación. The actualización de pesos equation for the original form of the algorithm can be written as

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \alpha \frac{\epsilon_k \mathbf{X}_k}{|\mathbf{X}_k|^2}. \quad (10)$$

The time index or ciclo de adaptación number is k . $\mathbf{W}_{k+1}$ is the next value of the vector de pesos, $\mathbf{W}_k$ is the present value of the vector de pesos, and $\mathbf{X}_k$ is the present vector de entrada de patrón. The present error lineal ϵ_k is defined to be the difference between the respuesta deseada d_k and the linear output $s_k = \mathbf{W}_k^T \mathbf{X}_k$ before adaptation:

$$\epsilon_k \triangleq d_k - \mathbf{W}_k^T \mathbf{X}_k. \quad (11)$$

Cambiar los pesos produce un cambio correspondiente en el error:

$$\Delta \epsilon_k = \Delta(d_k - \mathbf{W}_k^T \mathbf{X}_k) = -\mathbf{X}_k^T \Delta \mathbf{W}_k. \quad (12)$$

In accordance with the α -LMS rule of Eq. (10), the cambio de pesos is as follows:

$$\Delta \mathbf{W}_k = \mathbf{W}_{k+1} - \mathbf{W}_k = \alpha \frac{\epsilon_k \mathbf{X}_k}{|\mathbf{X}_k|^2}. \quad (13)$$

Combining Eqs. (12) and (13), we obtain

$$\Delta \epsilon_k = -\alpha \frac{\epsilon_k \mathbf{X}_k^T \mathbf{X}_k}{|\mathbf{X}_k|^2} = -\alpha \epsilon_k. \quad (14)$$

Therefore, the error is reduced by a factor of α as the weights are changed while holding the patrón de entrada fixed. Presenting a new patrón de entrada starts the next ciclo de adaptación. The next error is then reduced by a factor of α , and the process continues. The peso inicial vector is usually chosen to be zero and is adapted until convergence. In entornos no estacionarios, the weights are generally adapted continually.




Figure 12. Weight correction by the LMS rule.

The choice of α controls stability and speed of convergence [30]. For vector de entradas de patrón independent over time, stability is ensured for most practical purposes if

$$0 < \alpha < 2. \quad (15)$$

Making α greater than 1 generally does not make sense, since the error would be overcorrected. Total error correction comes with $\alpha = 1$. A practical range for α is

$$0.1 < \alpha < 1.0. \quad (16)$$

This algorithm is self-normalizing in the sense that the choice of α does not depend on the magnitude of the input signals. The actualización de pesos is collinear with the patrón de entrada and of a magnitude inversely proportional to $|\mathbf{X}_k|^2$. With binary ± 1 inputs, $|\mathbf{X}_k|^2$ is equal to the number of weights and does not vary from pattern to pattern. If the binary inputs are the usual 1 and 0, no adaptation occurs for weights with 0 inputs, while with ± 1 inputs, all weights are adapted each cycle and convergence tends to be faster. For this reason, the symmetric inputs $+1$ and -1 are generally preferred.

Figure 12 provides a geometrical picture of how the α -LMS rule works. In accord with Eq. (13), $\mathbf{W}_{k+1}$ equals $\mathbf{W}_k$ added to $\Delta \mathbf{W}_k$, and $\Delta \mathbf{W}_k$ is parallel with the vector de entrada de patrón $\mathbf{X}_k$. From Eq. (12), the change in error is equal to the negative producto punto of $\mathbf{X}_k$ and $\Delta \mathbf{W}_k$. Since the α -algorithm LMS selects $\Delta \mathbf{W}_k$ to be collinear with $\mathbf{X}_k$, the desired error correction is achieved with a weight change of the smallest possible magnitude. When adapting to respond properly to a new patrón de entrada, the responses to previous patrones de entrenamiento are therefore minimally disturbed, on the average.

The α -algorithm LMS corrects error, and if all patrones de entrada are all of equal length, it minimizes error cuadrático medio [30]. The algorithm is best known for this property.

4.2 Nonlinear Rules

The α -algorithm LMS is a linear rule that makes error corrections that are proportional to the error. It is known [105] that in some cases this linear rule may fail to separate patrones de entrenamiento that are linealmente separable. Where this creates difficulties, nonlinear rules may be used. In the next sections, we describe early nonlinear rules, which were devised by Rosenblatt [5, 106] and Mays [105]. These nonlinear rules also make vector de pesos changes collinear with the vector de




Figure 13. Rosenblatt's α -Perceptron.

entrada de patrón (the direction which causes minimal disturbance), changes that are based on the error lineal but are not directly proportional to it.

The Perceptron Learning Rule: The Rosenblatt α -Perceptron [5, 106], diagrammed in Fig. 13, processed patrones de entrada with a first layer of sparse randomly connected fixed dispositivos lógicos. The outputs of the fixed first layer fed a second layer, which consisted of a single adaptive linear threshold element. Other than the convention that its input signals were $\{1,0\}$ binary, and that no peso de sesgo was included, this element is equivalent to the Adaline element. The regla de aprendizaje for the α -Perceptron is very similar to LMS, but its behavior is in fact quite different.

It is interesting to note that Rosenblatt's regla de aprendizaje del Perceptron was first presented in 1960 [106], and Widrow and Hoff's LMS rule was first presented the same year, a few months later [59]. These rules were developed independently in 1959.

The elemento adaptativo de umbral of the α -Perceptron is shown in Fig. 14. Adapting with the regla del Perceptron makes use of the "quantizer error" $\tilde{\epsilon}_k$, defined to be the difference between the desired response and the output of the quantizer

$$\tilde{\epsilon}_k \triangleq d_k - y_k. \quad (17)$$

The regla del Perceptron, sometimes called the Perceptron convergence procedure, does not adapt the weights if the decisión de salida y_k is correct, that is, if $\tilde{\epsilon}_k = 0$. If the decisión de salida disagrees with the binary desired response d_k , however, adaptation is effected by adding the vector de entrada to the vector de pesos when the error $\tilde{\epsilon}_k$ is positive, or subtracting the vector de entrada from the vector de pesos when the error $\tilde{\epsilon}_k$ is negative. Thus, half the product of the vector de entrada and the error del cuantificador $\tilde{\epsilon}_k$ is added to the vector de pesos. The regla del Perceptron is identical to the α -algoritmo LMS, except that with the regla del Perceptron, half of the error del cuantificador $\tilde{\epsilon}_k/2$ is used in place of the normalized error lineal $\epsilon_k/|\mathbf{X}_k|^2$ of the α -LMS rule. The regla del Perceptron is nonlinear, in contrast to the LMS rule, which is linear (compare Figs. 2 and 14). Nonetheless, the regla del Perceptron can be written in a form very similar to the α -LMS rule of Eq. (10):

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \alpha \frac{\tilde{\epsilon}_k}{2} \mathbf{X}_k. \quad (18)$$




Figure 14. The elemento adaptativo de umbral of the Perceptron.

Rosenblatt normally set α to one. In contrast to α -LMS, the choice of α does not affect the stability of the algoritmo del Perceptron, and it affects tiempo de convergencia only if the peso inicial vector is nonzero. Also, while α -LMS can be used with either analog or binary respuestas deseadas, Rosenblatt's rule can be used only with binary respuestas deseadas.

The regla del Perceptron stops adapting when the patrones de entrenamiento are correctly separated. There is no restraining force controlling the magnitude of the weights, however. The direction of the vector de pesos, not its magnitude, determines the decision function. The regla del Perceptron has been proven to be capable of separating any linealmente separable set of training patterns [5, 46, 105, 107]. If the patrones de entrenamiento are not linealmente separable, the algoritmo del Perceptron goes on forever, and often does not yield a low-error solution, even if one exists. In most cases, if the conjunto de entrenamiento is not separable, the vector de pesos tends to gravitate toward zero¹² so that even if α is very small, each adaptation can dramatically affect the switching function implemented by the Perceptron.

This behavior is very different from that of the α -algoritmo LMS. Continued use of α -LMS does not lead to an unreasonable weight solution if the pattern set is not linealmente separable. Nor, however, is this algorithm guaranteed to separate any linealmente separable pattern set. α -LMS typically comes close to achieving such separation, but its objective is different—error reduction at the salida lineal of the adaptive element.

Rosenblatt also introduced variants of the fixed-increment rule that we have discussed thus far. A popular one was the absolute-correction version of the regla del Perceptron.¹³ This

¹²This results because the length of the vector de pesos decreases with each adaptation that does not cause the salida lineal s_k to change sign and assume a magnitude greater than that before adaptation. Although there are exceptions, for most problems this situation occurs only rarely if the vector de pesos is much longer than the weight increment vector.

¹³The terms "fixed-increment" and "absolute correction" are due to Nils-son [46]. Rosenblatt referred to methods of these types, respectively, as quantized and nonquantized reglas de aprendizaje.

rule is identical to that stated in Eq. (18) except the increment size α is chosen with each presentation to be the smallest integer which corrects the error de salida in one presentation. If the conjunto de entrenamiento is separable, this variant has all the characteristics of the fixed-increment version with α set to 1, except that it usually reaches a solution in fewer presentations.

Mays's Algorithms: In his tesis doctoral [105], Mays described an “increment adaptation” rule¹⁴ and a “modified relaxation adaptation” rule. The fixed-increment version of the regla del Perceptrón is a special case of the increment adaptation rule.

Increment adaptation in its general form involves the use of a “zona muerta” for the salida lineal s_k , equal to $\pm\gamma$ about zero. All respuestas deseadas are ± 1 (refer to Fig. 14). If the salida lineal s_k falls outside the zona muerta ($|s_k| \geq \gamma$), adaptation follows a normalized variant of the fixed-increment regla del Perceptrón (with $\alpha/|X_k|^2$ used in place of α). If the salida lineal falls within the zona muerta, whether or not the output response y_k is correct, the weights are adapted by the normalized variant of the regla del Perceptrón as though the respuesta de salida y_k had been incorrect. The weight update rule for Mays's increment algoritmo de adaptación can be written mathematically as

$$W_{k+1} = \begin{cases} W_k + \alpha \tilde{\epsilon}_k \frac{X_k}{2|X_k|^2} & \text{if } |s_k| \geq \gamma \\ W_k + \alpha d_k \frac{X_k}{|X_k|^2} & \text{if } |s_k| < \gamma \end{cases} \quad (19)$$

where $\tilde{\epsilon}_k$ is the error del cuantificador of Eq. (17).

With the zona muerta $\gamma = 0$, Mays's increment algoritmo de adaptación reduces to a normalized version of the regla del Perceptrón (18). Mays proved that if the training patterns are linealmente separable, increment adaptation will always converge and separate the patterns in a finite number of steps. He also showed that use of the zona muerta reduces sensitivity to weight errors. If the conjunto de entrenamiento is not linearly separable, Mays's increment adaptation rule typically performs much better than the regla del Perceptrón because a sufficiently large zona muerta tends to cause the weight vector to adapt away from zero when any reasonably good solution exists. In such cases, the vector de pesos may sometimes appear to meander rather aimlessly, but it will typically remain in a region associated with relatively low average error.

The increment adaptation rule changes the weights with increments that generally are not proportional to the error lineal ϵ_k . The other Mays rule, modified relaxation, is closer to α -LMS in its use of the error lineal ϵ_k (refer to Fig. 2). The respuesta deseada and the quantizer output levels are binary ± 1 . If the quantizer output y_k is wrong or if the salida lineal s_k falls within the zona muerta $\pm\gamma$, adaptation follows α -LMS to reduce the error lineal. If the quantizer output y_k is correct and the salida lineal s_k falls outside the zona muerta, the weights are not adapted. The actualización de pesos rule for this algorithm can be written as

$$W_{k+1} = \begin{cases} W_k & \text{if } \tilde{\epsilon}_k = 0 \text{ and } |s_k| \geq \gamma \\ W_k + \alpha \epsilon_k \frac{X_k}{|X_k|^2} & \text{otherwise} \end{cases} \quad (20)$$

¹⁴The increment adaptation rule was proposed by others before Mays, though from a different perspective [107].




Figure 15. A five-Adaline example of the Madaline I architecture.

where $\tilde{\epsilon}_k$ is the error del cuantificador of Eq. (17).

If the zona muerta γ is set to ∞ , this algorithm reduces to the α -algoritmo LMS (10). Mays showed that, for zona muerta $0 < \gamma < 1$ and tasa de aprendizaje $0 < \alpha \leq 2$, this algorithm will converge and separate any linealmente separable input set in a finite number of steps. If the conjunto de entrenamiento is not linealmente separable, this algorithm performs much like Mays's increment adaptation rule.

Los dos algoritmos de Mays logran resultados similares de separación de patrones. The choice of α does not affect stability, although it does affect tiempo de convergencia. The two rules differ in their convergence properties but there is no consensus on which is the better algorithm. Algorithms like these can be quite useful, and we believe that there are many more to be invented and analyzed.

The α -algoritmo LMS, the Perceptrón procedure, and Mays's algorithms can all be used for adapting the single Adaline element or they can be incorporated into procedures for adapting networks of such elements. Multilayer network adaptation procedures that use some of these algorithms are discussed in the following.

5 Error-Correction Rules—Multi-Element Networks

The algorithms discussed next are the Widrow–Hoff Madaline rule from the early 1960s, now called Regla Madaline I (MRI), and Regla Madaline II (MRII), developed by Widrow and Winter in 1987.

5.1 Regla Madaline I

The MRI rule allows the adaptation of a first layer of con límite estricto (signum) Adaline elements whose outputs provide inputs to a second layer, consisting of a single fixed-threshold-logic element which may be, for example, the compuerta OR, AND gate, or tomador de voto mayoritario discussed previously. The weights of the Adalines are initially set to small random values.

Figure 15 shows a Madaline I architecture with five completamente conectada de la primera capa Adalines. La segunda capa es un elemento mayoritario (MAJ). Because the de la segunda capa logic element is fixed and known, it is possible to determine which de la primera capa Adalines can be adapted to correct an error de salida. The Adalines in the first layer assist each other in solving problems by automatic distribución de carga.

One procedure for training the network in Fig. 15 follows. A pattern is presented, and if the respuesta de salida de the elemento mayoritario matches the desired response, no adaptation takes place. However, if, for instance, the desired response is $+1$ and three of the five Adalines read -1 for a given input pattern, one of the latter three must be adapted to the $+1$ state. The element that is adapted by MRI is the one whose salida lineal s_k is closest to zero—the one whose analog response is closest to the respuesta deseada. If more of the Adalines were originally in the -1 state, enough of them are adapted to the $+1$ state to make the majority decision equal $+1$. The elements adapted are those whose salida lineales are closest to zero. A similar procedure is followed when the respuesta deseada is -1 . When adapting a given element, the weight vector can be moved in the LMS direction far enough to reverse the Adaline's output (absolute correction, or "fast" learning), or it can be adapted by the small increment determined by the α -algoritmo LMS (statistical, or "slow" learning). The one respuesta deseada d_k is used for all Adalines that are adapted. The procedure can also be modified to allow one of Mays's rules to be used. In that event, for the case we have considered (majority output element), adaptations take place if at least half of the Adalines either have outputs differing from the respuesta deseada or have analog outputs which are in the dead zone. By setting the zona muerta of Mays's increment adaptation rule to zero, the weights can also be adapted by Rosenblatt's regla del Perceptrón.

Differences in initial conditions and the results of subsequent adaptation cause the various elements to take "responsibility" for certain parts of the training problem. El principio básico de distribución de carga se resume así: *Assign responsibility to the Adaline or Adalines that can most easily assume it.*

In Fig. 15, the "asignador de tareas," a purely mechanized process, assigns responsibility during training by transferring the appropriate adapt commands and respuesta deseada signals to the selected Adalines. The asignador de tareas utilizes linear-output information. Load sharing is important, since it results in the various adaptive elements developing individual vectores de pesos. If all the weight vectors were the same, there would be no point in having more than one element in the first layer.

Al entrenar la Madaline, la secuencia de presentación de patrones debe ser aleatoria. Experimenting with this, Ridgway [76] found that cyclic presentation of the patterns could lead to cycles of adaptation. These cycles would cause the weights of the entire Madaline to cycle, preventing convergence.

The adaptive system of Fig. 15 was suggested by common sense, and was found to work well in simulations. Ridgway found that the probability that a given Adaline will be adapted in response to an patrón de entrada is greatest if that element had taken such responsibility during the previous adapt cycle when the pattern was most recently presented. The division of responsibility stabilizes at the same time that the responses of individual elements stabilize to their share of the load. When




Figure 16. Typical two-layer Madaline II architecture.

the training problem is not perfectly separable by this system, the adaptation process tends to minimize probabilidad de error, although it is possible for the algorithm to "hang up" on óptimos locales.

The Madaline structure of Fig. 15 has 2 layers—the first layer consists of adaptive logic elements, the second of fixed logic. A variety of fixed-logic devices could be used for the second layer. A variety of MRI adaptation rules were devised by Hoff [75] that can be used with all possible fixed-logic output elements. An easily described training procedure results when the output element is an compuerta OR. During training, if the desired output for a given patrón de entrada is $+1$, only the one Adaline whose salida lineal is closest to zero would be adapted if any adaptation is needed—in other words, if all Adalines give -1 outputs. If the desired output is -1 , all elements must give -1 outputs, and any giving $+1$ outputs must be adapted.

La regla MRI obedece el "principio de mínima perturbación" en el siguiente sentido. No more Adaline elements are adapted than necessary to correct the decisión de salida and any dead-zone constraint. The elements whose salida lineales are nearest to zero are adapted because they require the smallest cambios de pesos to reverse their output responses. Furthermore, whenever an Adaline is adapted, the weights are changed in the direction of its vector de entrada, providing the requisite error correction with minimal cambio de pesos.

5.2 Regla Madaline II

The MRI rule was recently extended to allow the adaptation of multilayer binary networks by Winter and Widrow with the introduction of Regla Madaline II (MRII) [43, 83, 108]. A typical two-layer MRII network is shown in Fig. 16. Los pesos en ambas capas son adaptativos.

El entrenamiento con la regla MRII es similar al entrenamiento con el algoritmo MRI. The weights are initially set to small random values. Training patterns are presented in a random sequence. If the network produces an error during a training presentation, we begin by adapting de la primera capa Adalines.

By the principio de mínima perturbación, we select the de la primera capa Adaline with the smallest salida lineal magnitude and perform a “adaptación de prueba” by inverting its binary output. This can be done without adaptation by adding a perturbation Δs of suitable amplitude and polarity to the Adaline’s sum (refer to Fig. 16). If the error de salida de Hamming is reduced by this bit inversion, that is, if the number of errores de salida is reduced, the perturbation Δs is removed and the weights of the selected Adaline element are changed by α -LMS in a direction collinear with the corresponding vector de entrada—the direction that reinforces the inversión de bit with minimal disturbance to the weights. Conversely, if the adaptación de prueba does not improve the network response, no weight adaptation is performed.

After finishing with the first element, we perturb and update other Adalines in the first layer which have “sufficiently small” linear-output magnitudes. Further error reductions can be achieved, if desired, by reversing pairs, triples, and so on, up to some predetermined limit. After exhausting possibilities with the first layer, we move on to the next layer and proceed in a like manner. When the final layer is reached, each of the output elements is adapted by α -LMS. At this point, a new patrón de entrenamiento is selected at random and the procedure is repeated. The goal is to reduce error de Hamming with each presentation, thereby hopefully minimizing the average error de Hamming over the conjunto de entrenamiento. Like MRI, the procedure can be modified so that adaptations follow an absolute correction rule or one of Mays’s rules rather than α -LMS. Like MRI, MRII can “hang up” on óptimos locales.

6 Steepest-Descent Rules—Single Threshold Element

Thus far, we have described a variety of adaptation rules that act to reduce error with the presentation of each patrón de entrenamiento. Often, the objective of adaptation is to reduce error averaged in some way over the conjunto de entrenamiento. The most common error function is error cuadrático medio (MSE), although in some situations other error criteria may be more appropriate [109–111]. The most popular approaches to MSE reduction in both mono-elemento and redes multi-elemento are based upon the method of descenso más pronunciado. More sophisticated gradient approaches such as cuasi-Newton [30, 112, 113] and conjugate gradient [113, 114] techniques often have better convergence properties, but the conditions under which the additional complexity is warranted are not generally known. The discussion that follows is restricted to minimization of MSE by the method of descenso más pronunciado [116, 117]. More sophisticated learning procedures usually require many of the same computations used in the basic descenso más pronunciado procedure.

Adaptation of a network by descenso más pronunciado starts with an arbitrary initial value $\mathbf{W}_0$ for the system’s vector de pesos. The gradient of the MSE function is measured and the vector de pesos is altered in the direction corresponding to the negative of the measured gradient. This procedure is repeated, causing the MSE to be successively reduced on average and causing

the vector de pesos to approach a locally optimal value.

El método de descenso más pronunciado puede describirse mediante la relación

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \mu(-\nabla_k) \quad (21)$$

where μ is a parameter that controls stability and rate of convergence, and ∇_k is the value of the gradient at a point on the superficie ECM corresponding to $\mathbf{W} = \mathbf{W}_k$.

To begin, we derive rules for descenso más pronunciado minimization of the MSE associated with a single Adaline element. These rules are then generalized to apply to full-blown neural networks. Like reglas de corrección de error, the most practical and efficient descenso más pronunciado rules typically work with one pattern at a time. They minimize error cuadrático medio, approximately, averaged over the entire set of training patterns.

6.1 Linear Rules

Steepest-descent rules for the single threshold element are said to be linear if cambios de pesos are proportional to the error lineal, the difference between the respuesta deseada d_k and the salida lineal of the element s_k .

Mean-Square Error Surface of the Linear Combiner: In this section we demonstrate that the superficie ECM of the combinador lineal of Fig. 1 is a función cuadrática of the weights, and thus easily traversed by descenso por gradiente.

Let the patrón de entrada $\mathbf{X}_k$ and the associated respuesta deseada d_k be drawn from a statistically población estacionaria. During adaptation, the vector de pesos varies so that even with stationary inputs, the output s_k and error ϵ_k will generally be nonstationary. Care must be taken in defining the MSE since it is time-varying. La única posibilidad es un promedio de conjunto, definido a continuación.

At the k th iteration, let the vector de pesos be $\mathbf{W}_k$. Squaring and expanding Eq. (11) yields

$$\epsilon_k^2 = (d_k - \mathbf{X}_k^T \mathbf{W}_k)^2 \quad (22)$$

$$= d_k^2 - 2d_k \mathbf{X}_k^T \mathbf{W}_k + \mathbf{W}_k^T \mathbf{X}_k \mathbf{X}_k^T \mathbf{W}_k. \quad (23)$$

Now assume an ensemble of identical combinador lineal adaptativos, each having the same vector de pesos $\mathbf{W}_k$ at the k th iteration. Let each combiner have individual inputs $\mathbf{X}_k$ and d_k derived from ergódico estacionario ensembles. Each combiner will produce an individual error ϵ_k represented by Eq. (23). Averaging Eq. (23) over the ensemble yields

$$E[\epsilon_k^2]_{\mathbf{W}=\mathbf{W}_k} = E[d_k^2] - 2E[d_k \mathbf{X}_k^T] \mathbf{W}_k + \mathbf{W}_k^T E[\mathbf{X}_k \mathbf{X}_k^T] \mathbf{W}_k. \quad (24)$$

Defining the vector $\mathbf{P}$ as the correlación cruzada between the respuesta deseada (a scalar) and the $\mathbf{X}$ -vector¹⁵ then yields

$$\mathbf{P}^T \triangleq E[d_k \mathbf{X}_k^T] = E[d_k, d_k x_{1k}, \dots, d_k x_{nk}]^T. \quad (25)$$

The matriz de correlación de entrada $\mathbf{R}$ is defined in terms of

¹⁵We assume here that $\mathbf{X}$ includes a bias component $x_{0k} = +1$.

Note that the MSE is a quadratic function of the weights. It is a convex hyperparaboloidal surface, a function that never goes negative. Figure 17 shows a typical MSE surface




Figure 17. Typical mean-square-error surface of a combinatorial lineal.

the promedio de conjunto

$$\mathbf{R} \triangleq E[\mathbf{X}_k \mathbf{X}_k^T] = E \begin{bmatrix} 1 & x_{1k} & \cdots & x_{nk} \\ x_{1k} & x_{1k}x_{1k} & \cdots & x_{1k}x_{nk} \\ \vdots & \vdots & \ddots & \vdots \\ x_{nk} & x_{nk}x_{1k} & \cdots & x_{nk}x_{nk} \end{bmatrix}. \quad (26)$$

This matrix is real, symmetric, and definida positiva, or in rare cases, positive semi-definite. The MSE ξ_k can thus be expressed as

$$\xi_k \triangleq E[\epsilon_k^2]_{\mathbf{W}=\mathbf{W}_k} = E[d_k^2] - 2\mathbf{P}^T \mathbf{W}_k + \mathbf{W}_k^T \mathbf{R} \mathbf{W}_k. \quad (27)$$

Nótese que el ECM es una función cuadrática de los pesos. It is a convex hyperparaboloidal surface, a function that never goes negative. Figure 17 shows a typical superficie ECM for a combinatorial lineal with two weights. The position of a point on the grid in this figure represents the value of the Adaline's two weights. The height of the surface at each point represents MSE over the training set when the Adaline's weights are fixed at the values associated with the grid point. Adjusting the weights involves descending along this surface toward the unique minimum point ("the bottom of the bowl") by the method of steepest descent.

The gradient ∇_k of the MSE function with $\mathbf{W} = \mathbf{W}_k$ is obtained by differentiating Eq. (27):

$$\nabla_k \triangleq \begin{bmatrix} \frac{\partial E[\epsilon_k^2]}{\partial w_{0k}} \\ \vdots \\ \frac{\partial E[\epsilon_k^2]}{\partial w_{nk}} \end{bmatrix}_{\mathbf{W}=\mathbf{W}_k} = -2\mathbf{P} + 2\mathbf{R}\mathbf{W}_k. \quad (28)$$

Esta es una función lineal de los pesos. The optimal vector de pesos $\mathbf{W}^*$, generally called the vector de pesos de Wiener, is obtained from Eq. (28) by setting the gradient to zero:

$$\mathbf{W}^* = \mathbf{R}^{-1} \mathbf{P}. \quad (29)$$

Esta es una forma matricial de la ecuación de Wiener-Hopf [118–120]. In the next section we examine μ -LMS, an

algorithm which enables us to obtain an accurate estimate of $\mathbf{W}^*$ without first computing $\mathbf{R}^{-1}$ and $\mathbf{P}$.

The μ -LMS Algorithm: The μ -algoritmo LMS works by performing approximate descenso más pronunciado on the superficie ECM in espacio de pesos. Because it is a función cuadrática of the weights, this surface is convex and has a unique (global) minimum.¹⁶ An gradiente instantáneo based upon the square of the instantaneous error lineal is

$$\hat{\nabla}_k = \frac{\partial \epsilon_k^2}{\partial \mathbf{W}_k} = \begin{bmatrix} \frac{\partial \epsilon_k^2}{\partial w_{0k}} \\ \vdots \\ \frac{\partial \epsilon_k^2}{\partial w_{nk}} \end{bmatrix}. \quad (30)$$

LMS works by using this crude estimación del gradiente in place of the gradiente verdadero ∇_k of Eq. (28). Making this replacement into Eq. (21) yields

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \mu(-\hat{\nabla}_k) = \mathbf{W}_k - \mu \frac{\partial \epsilon_k^2}{\partial \mathbf{W}_k}. \quad (31)$$

The gradiente instantáneo is used because it is readily available from a single data sample. El gradiente verdadero es generalmente difícil de obtener. Computing it would involve averaging the gradiente instantáneos associated with all patterns in the conjunto de entrenamiento. Esto es usualmente impracticable y casi siempre ineficiente.

Performing the differentiation in Eq. (31) and replacing the error lineal by definition (11) gives

$$\mathbf{W}_{k+1} = \mathbf{W}_k - 2\mu\epsilon_k \frac{\partial \epsilon_k}{\partial \mathbf{W}_k} = \mathbf{W}_k - 2\mu\epsilon_k \frac{\partial (d_k - \mathbf{W}_k^T \mathbf{X}_k)}{\partial \mathbf{W}_k}. \quad (32)$$

Noting that d_k and $\mathbf{X}_k$ are independent of $\mathbf{W}_k$ yields

$$\mathbf{W}_{k+1} = \mathbf{W}_k + 2\mu\epsilon_k \mathbf{X}_k. \quad (33)$$

This is the μ -algoritmo LMS. The constante de aprendizaje μ determines stability and tasa de convergencia. For patrones de entrada independent over time, convergence of the mean and variance of the vector de pesos is ensured [30] for most practical purposes if

$$0 < \mu < \frac{1}{\text{trace}[\mathbf{R}]} \quad (34)$$

where $\text{trace}[\mathbf{R}] = \sum(\text{diagonal elements of } \mathbf{R})$ is the average signal power of the $\mathbf{X}$ -vectors, that is, $E[\mathbf{X}^T \mathbf{X}]$. With μ set within this range,¹⁷ the μ -algoritmo LMS converges in the mean to $\mathbf{W}^*$, the optimal Wiener solution discussed above. A proof of this can be found in [30].

In the μ -algoritmo LMS, and other iterative descenso más pronunciado procedures, use of the gradiente instantáneo is perfectly justified if the step size is small. For small μ , $\mathbf{W}$ will

¹⁶If the matriz de autocorrelación of the vector de patrón set has m zero valores propios, the minimum MSE solution will be an m -dimensional subspace in weight space [30].

¹⁷Horowitz and Senne [121] have proven that (34) is not sufficient in general to guarantee convergence of the vector de pesos's variance. For input patterns generated by a zero-mean Gaussian process independent over time, instability can occur in the worst case if μ is greater than $1/(3 \text{ trace}[\mathbf{R}])$.

remain essentially constant over a relatively small number of training presentations K . The total cambio de pesos during this period will be proportional to

$$\begin{aligned} -\sum_{i=0}^{K-1} \frac{\partial \epsilon_{k+i}^2}{\partial \mathbf{W}_{k+i}} &\simeq -K \left(\frac{1}{K} \sum_{i=0}^{K-1} \frac{\partial \epsilon_{k+i}^2}{\partial \mathbf{W}_k} \right) \\ &= -K \frac{\partial}{\partial \mathbf{W}_k} \left(\frac{1}{K} \sum_{i=0}^{K-1} \epsilon_{k+i}^2 \right) \simeq -K \frac{\partial \xi}{\partial \mathbf{W}_k} \end{aligned} \quad (35)$$

where ξ denotes the MSE function. Thus, on average the weights follow the true gradient. It is shown in [30] that the gradiente instantáneo is an unbiased estimate of the gradiente verdadero.

Comparison of μ -LMS and α -LMS: We have now presented two forms of the algoritmo LMS, α -LMS (10) in Section IV-A and μ -LMS (33) in the last section. They are very similar algorithms, both using the LMS instantaneous gradient. α -LMS is self-normalizing, with the parameter α determining the fraction of the instantaneous error to be corrected with each adaptation. μ -LMS is a constant-coefficient linear algorithm which is considerably easier to analyze than α -LMS. Comparing the two, the α -algoritmo LMS is like the μ -algoritmo LMS with a continually variable constante de aprendizaje. Although α -LMS is somewhat more difficult to implement and analyze, it has been demonstrated experimentally to be a better algorithm than μ -LMS when the valores propios of the input matriz de autocorrelación $\mathbf{R}$ are highly disparate, giving faster convergence for a given level of ruido de gradiente¹⁸ propagated into the weights. It will be shown next that μ -LMS has the advantage that it will always converge in the mean to the minimum MSE solution, while α -LMS may converge to a somewhat biased solution.

We begin with α -LMS of Eq. (10):

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \alpha \frac{\epsilon_k \mathbf{X}_k}{|\mathbf{X}_k|^2}. \quad (36)$$

Reemplazando el error con su definición (11) y reordenando términos se obtiene

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \alpha \frac{(d_k - \mathbf{W}_k^T \mathbf{X}_k) \mathbf{X}_k}{|\mathbf{X}_k|^2} \quad (37)$$

$$= \mathbf{W}_k + \alpha \left(\frac{d_k}{|\mathbf{X}_k|} - \mathbf{W}_k^T \frac{\mathbf{X}_k}{|\mathbf{X}_k|} \right) \frac{\mathbf{X}_k}{|\mathbf{X}_k|}. \quad (38)$$

We define a new conjunto de entrenamiento of vectores de patrón and respuestas deseadas $\{\bar{\mathbf{X}}_k, \bar{d}_k\}$ by normalizing elements of the original conjunto de entrenamiento as follows,¹⁹

$$\bar{\mathbf{X}}_k \triangleq \frac{\mathbf{X}_k}{|\mathbf{X}_k|}, \quad \bar{d}_k \triangleq \frac{d_k}{|\mathbf{X}_k|}. \quad (39)$$

Eq. (38) then becomes

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \alpha (\bar{d}_k - \mathbf{W}_k^T \bar{\mathbf{X}}_k) \bar{\mathbf{X}}_k. \quad (40)$$

¹⁸Gradient noise is the difference between the estimación del gradiente and the gradiente verdadero.

¹⁹The idea of a normalized conjunto de entrenamiento was suggested by Derick Nguyen.

This is the μ -LMS rule of Eq. (33) with 2μ replaced by α . The weight adaptations chosen by the α -LMS rule are equivalent to those of the μ -algoritmo LMS presented with a different conjunto de entrenamiento—the normalized conjunto de entrenamiento defined by (39). The solution that will be reached by the μ -LMS algorithm is the solución de Wiener of this conjunto de entrenamiento

$$\bar{\mathbf{W}}^* = (\bar{\mathbf{R}})^{-1} \bar{\mathbf{P}} \quad (41)$$

where

$$\bar{\mathbf{R}} = E[\bar{\mathbf{X}}_k \bar{\mathbf{X}}_k^T] \quad (42)$$

is the matriz de correlación de entrada of the normalized conjunto de entrenamiento and the vector

$$\bar{\mathbf{P}} = E[\bar{d}_k \bar{\mathbf{X}}_k] \quad (43)$$

is the correlación cruzada between the normalized input and the normalized desired response. Therefore α -LMS converges in the mean to the solución de Wiener of the normalized conjunto de entrenamiento. When the vectores de entrada are binary with ± 1 components, all vectores de entrada have the same magnitude and the two algorithms are equivalent. For nonbinary patrones de entrenamiento, however, the solución de Wiener of the normalized conjunto de entrenamiento generally is no longer equal to that of the original problem, so α -LMS converges in the mean to a somewhat biased version of the optimal least-squares solution.

The idea of a normalized conjunto de entrenamiento can also be used to relate the stable ranges for the constante de aprendizajes α and μ in the two algorithms. The stable range for α in the α -algoritmo LMS given in Eq. (15) can be computed from the corresponding range for μ given in Eq. (34) by replacing $\mathbf{R}$ and μ in Eq. (34) by $\bar{\mathbf{R}}$ and $\alpha/2$, respectively, and then noting that $\text{trace}[\bar{\mathbf{R}}]$ is equal to one:

$$0 < \alpha < \frac{2}{\text{trace}[\bar{\mathbf{R}}]}, \quad \text{or} \quad 0 < \alpha < 2. \quad (44)$$

6.2 Nonlinear Rules

The Adaline elements considered thus far use at their outputs either de límite estricto quantizers (signums), or no nonlinearity at all. The input–output mapping of the cuantificador de límite estricto is $y_k = \text{sgn}(s_k)$. Other forms of nonlinearity have come into use in the past two decades, primarily of the sigmoide type. These nonlinearities provide saturation for toma de decisiones, yet they have differentiable input–output characteristics that facilitate adaptivity. We generalize the definition of the Adaline element to include the possible use of a sigmoide in place of the signum, and then determine suitable algoritmo de adaptaci3n.

Fig. 18 shows a “Adaline sigmoideal” element which incorporates a sigmoideal nonlinearity. The input–output relation of the sigmoide can be denoted by $y_k = \text{sgm}(s_k)$. Una funci3n sigmoidee típica es la tangente hiperb3lica:

$$y_k = \tanh(s_k) = \left(\frac{1 - e^{-2s_k}}{1 + e^{-2s_k}} \right). \quad (45)$$

We shall adapt this Adaline with the objective of minimizing the mean square of the sigmoide error $\tilde{\epsilon}_k$, defined as

$$\tilde{\epsilon}_k \triangleq d_k - y_k = d_k - \text{sgm}(s_k). \quad (46)$$

defined as

$$\tilde{\epsilon}_k \triangleq d_k - y_k = d_k - \text{sgm}(s_k). \quad (46)$$

Backpropagation for the Sigmoid Adaline: Our objective is to minimize $E[(\tilde{\epsilon}_k)^2]$, averaged over the set of training patterns, by proper choice of the weight vector. To accomplish this, we shall derive a backpropagation algorithm for the sigmoid Adaline element. An instantaneous gradient is obtained with each input vector presentation, and the method of steepest descent is used to minimize error as was done with the μ -LMS algorithm of Eq. (33).

Referring to Fig. 18, the instantaneous gradient estimate




Figure 18. Adaline with sigmoideal nonlinearity.

Retropropagación for the Sigmoid Adaline: Our objective is to minimize $E[(\tilde{\epsilon}_k)^2]$, averaged over the set of patrones de entrenamiento, by proper choice of the vector de pesos. To accomplish this, we shall derive a retropropagación algorithm for the Adaline sigmoideal element. An gradiente instantáneo is obtained with each vector de entrada presentation, and the method of descenso más pronunciado is used to minimize error as was done with the μ -algoritmo LMS of Eq. (33).

Referring to Fig. 18, the gradiente instantáneo estimate obtained during presentation of the k th vector de entrada $\mathbf{X}_k$ is given by

$$\hat{\nabla}_k = \frac{\partial(\tilde{\epsilon}_k)^2}{\partial \mathbf{W}_k} = 2\tilde{\epsilon}_k \frac{\partial \tilde{\epsilon}_k}{\partial \mathbf{W}_k}. \quad (47)$$

Differentiating Eq. (46) yields

$$\frac{\partial \tilde{\epsilon}_k}{\partial \mathbf{W}_k} = -\frac{\partial \text{sgm}(s_k)}{\partial \mathbf{W}_k} = -\text{sgm}'(s_k) \frac{\partial s_k}{\partial \mathbf{W}_k}. \quad (48)$$

We may note that

$$s_k = \mathbf{X}_k^T \mathbf{W}_k. \quad (49)$$

Therefore,

$$\frac{\partial s_k}{\partial \mathbf{W}_k} = \mathbf{X}_k. \quad (50)$$

Substituting into Eq. (48) gives

$$\frac{\partial \tilde{\epsilon}_k}{\partial \mathbf{W}_k} = -\text{sgm}'(s_k) \mathbf{X}_k. \quad (51)$$

Inserting this into Eq. (47) yields

$$\hat{\nabla}_k = -2\tilde{\epsilon}_k \text{sgm}'(s_k) \mathbf{X}_k. \quad (52)$$

Using this estimación del gradiente with the method of descenso más pronunciado provides a means for minimizing the error




Figure 19. Implementation of retropropagación for the Adaline sigmoideal element.

cuadrático medio even after the summed signal s_k goes through the nonlinear sigmoide. The algorithm is

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \mu(-\hat{\nabla}_k) \quad (53)$$

$$= \mathbf{W}_k + 2\mu\tilde{\epsilon}_k \text{sgm}'(s_k) \mathbf{X}_k. \quad (54)$$

El algoritmo (54) es el algoritmo de retropropagación para el elemento Adaline sigmoideal. The retropropagación name makes more sense when the algorithm is utilized in a red en capas, which will be studied below. Implementation of algorithm (54) is illustrated in Fig. 19.

If the sigmoide is chosen to be the tangente hiperbólica function (45), then the derivative $\text{sgm}'(s_k)$ is given by

$$\text{sgm}'(s_k) = \frac{\partial(\tanh(s_k))}{\partial s_k} = 1 - (\tanh(s_k))^2 = 1 - y_k^2. \quad (55)$$

Accordingly, Eq. (54) becomes

$$\mathbf{W}_{k+1} = \mathbf{W}_k + 2\mu\tilde{\epsilon}_k(1 - y_k^2) \mathbf{X}_k. \quad (56)$$

Regla Madaline III for the Sigmoid Adaline: The implementation of algorithm (54) (Fig. 19) requires accurate realization of the sigmoide function and its derivative function. These functions may not be realized accurately when implemented with hardware analógico. Indeed, in an analog network, each Adaline will have its own individual nonlinearities. Difficulties in adaptation have been encountered in practice with the algoritmo de retropropagación because of imperfections in the nonlinear functions.

To circumvent these problems a new algorithm has been devised by David Andes for adapting networks of Adaline sigmoideals. This is the Regla Madaline III (MRIII) algorithm. The idea of MRIII for a Adaline sigmoideal is illustrated in Fig. 20. La derivada de la función sigmoidee no se usa aquí. Instead, a small perturbation signal Δs is added to the sum s_k , and the effect of this perturbation upon output y_k and error $\tilde{\epsilon}_k$ is noted.

An instantaneous estimated gradient can be obtained as follows:

$$\hat{\nabla}_k = \frac{\partial(\tilde{\epsilon}_k)^2}{\partial \mathbf{W}_k} = \frac{\partial(\tilde{\epsilon}_k)^2}{\partial s_k} \frac{\partial s_k}{\partial \mathbf{W}_k} = \frac{\partial(\tilde{\epsilon}_k)^2}{\partial s_k} \mathbf{X}_k. \quad (57)$$

ered network, which will be studied below. Implementation of algorithm (54) is illustrated in Fig. 19.

If the sigmoid is chosen to be the hyperbolic tangent function (45), then the derivative $\text{sgm}'(s_k)$ is given by

$$\begin{aligned}\text{sgm}'(s_k) &= \frac{\partial(\tanh(s_k))}{\partial s_k} \\ &= 1 - (\tanh(s_k))^2 = 1 - y_k^2.\end{aligned}\quad (55)$$

Accordingly, Eq. (54) becomes

$$\mathbf{W}_{k+1} = \mathbf{W}_k + 2\mu \tilde{\epsilon}_k (1 - y_k^2) \mathbf{X}_k. \quad (56)$$

Madaline Rule III for the Sigmoid Adaline: The implementation of algorithm (54) (Fig. 19) requires accurate realization of the sigmoid function and its derivative function. These functions may not be realized accurately when implemented with analog hardware. Indeed, in an analog network, each Adaline will have its own individual nonlinearities. Difficulties in adaptation have been encountered in practice with the backpropagation algorithm because of imperfections in the nonlinear functions.

To circumvent these problems a new algorithm has been devised by David Andes for adapting networks of sigmoid Adalines. This is the Madaline Rule III (MRIII) algorithm.

The idea of MRIII for a sigmoid Adaline is illustrated in

Figure 20. Implementation of the MRIII algorithm for the Adaline sigmoideal element.

Since Δs is small,

$$\hat{\nabla}_k \simeq \left(\frac{\Delta(\tilde{\epsilon}_k)^2}{\Delta s} \right) \mathbf{X}_k. \quad (58)$$

Another way to obtain an approximate gradiente instantáneo by measuring the effects of the perturbation Δs can be obtained from Eq. (57).

$$\hat{\nabla}_k = \frac{\partial(\tilde{\epsilon}_k)^2}{\partial s_k} \mathbf{X}_k = 2\tilde{\epsilon}_k \frac{\partial \tilde{\epsilon}_k}{\partial s_k} \mathbf{X}_k \simeq 2\tilde{\epsilon}_k \left(\frac{\Delta \tilde{\epsilon}_k}{\Delta s} \right) \mathbf{X}_k. \quad (59)$$

En consecuencia, hay dos formas del algoritmo MRIII para la Adaline sigmoideal. They are based on the method of descenso más pronunciado, using the estimated instantaneous gradients:

$$\mathbf{W}_{k+1} = \mathbf{W}_k - \mu \left(\frac{\Delta(\tilde{\epsilon}_k)^2}{\Delta s} \right) \mathbf{X}_k \quad (60)$$

or,

$$\mathbf{W}_{k+1} = \mathbf{W}_k - 2\mu \tilde{\epsilon}_k \left(\frac{\Delta \tilde{\epsilon}_k}{\Delta s} \right) \mathbf{X}_k. \quad (61)$$

Para perturbaciones pequeñas, estas dos formas son esencialmente idénticas. Neither one requires *a priori* knowledge of the sigmoideal's derivative, and both are robust with respect to natural variations, biases, and drift in the analog hardware. Qué forma usar es una cuestión de conveniencia de implementación. The algorithm of Eq. (60) is illustrated in Fig. 20.

Regarding algorithm (61), some changes can be made to establish a point of interest. Note that, in accord with Eq. (46),

$$\tilde{\epsilon}_k = d_k - y_k. \quad (62)$$




Figure 21. The linear, sigmoideal, and signum errors of the Adaline.

Adding the perturbation Δs causes a change in $\tilde{\epsilon}_k$ equal to

$$\Delta \tilde{\epsilon}_k = -\Delta y_k. \quad (63)$$

Now, Eq. (61) may be rewritten as

$$\mathbf{W}_{k+1} = \mathbf{W}_k + 2\mu \tilde{\epsilon}_k \left(\frac{\Delta y_k}{\Delta s} \right) \mathbf{X}_k. \quad (64)$$

Since Δs is small, the ratio of increments may be replaced by a ratio of differentials, finally giving

$$\mathbf{W}_{k+1} \simeq \mathbf{W}_k + 2\mu \tilde{\epsilon}_k \frac{\partial y_k}{\partial s_k} \mathbf{X}_k \quad (65)$$

$$= \mathbf{W}_k + 2\mu \tilde{\epsilon}_k \text{sgm}'(s_k) \mathbf{X}_k. \quad (66)$$

Esto es idéntico al algoritmo de retropropagación (54) para la Adaline sigmoideal. Thus, retropropagación and MRIII are mathematically equivalent if the perturbation Δs is small, but MRIII is robust, even with analog implementations.

MSE Surfaces of the Adaline: Fig. 21 shows a combinatorial lineal connected to both sigmoideal and signum devices. Three errors, ϵ , $\tilde{\epsilon}_k$, and $\tilde{\tilde{\epsilon}}$ are designated in this figure. They are:

$$\begin{aligned}\text{error lineal} &= \epsilon = d - s \\ \text{sigmoideal error} &= \tilde{\epsilon} = d - \text{sgm}(s) \\ \text{signum error} &= \tilde{\tilde{\epsilon}} = d - \text{sgn}(\text{sgm}(s)) \\ &= d - \text{sgn}(s).\end{aligned}\quad (67)$$

To demonstrate the nature of the error cuadrático surfaces associated with these three types of error, a simple experiment with a two-input Adaline was performed. The Adaline was driven by a typical set of patrones de entrada and their associated binary $\{+1, -1\}$ respuestas deseadas. The función sigmoideal used was the hyperbolic tangent. The weights could have been adapted to minimize the error cuadrático medio of ϵ , $\tilde{\epsilon}$, or $\tilde{\tilde{\epsilon}}$. The superficies ECM of $E[(\epsilon)^2]$, $E[(\tilde{\epsilon})^2]$, $E[(\tilde{\tilde{\epsilon}})^2]$ plotted as functions of the two weight values, are shown in Figs. 22, 23, and 24, respectively.

Although the above experiment is not all encompassing, we can infer from it that minimizing the mean square of the error lineal is easy and minimizing the mean square of the sigmoideal




Figure 22. Example superficie ECM of error lineal.




Figure 23. Example superficie ECM of sigmoide error.

error is more difficult, but typically much easier than minimizing the mean square of the signum error. Only the error lineal is guaranteed to have an superficie ECM with a unique mínimo global (assuming invertible $\mathbf{R}$ -matrix). Las otras superficies ECM pueden tener óptimos locales [122, 123].

In nonlinear neural networks, gradient methods generally work better with sigmoide rather than signum nonlinearities. Smooth nonlinearities are required by the MRIII and técnica de retropropagación. Moreover, redes sigmoideales are capable of forming representaciones internas that are more complex than simple binary codes and, thus, these networks can often form decision regions that are more sophisticated than those associated with similar redes signum. In fact, if a noiseless infinite-precision Adaline sigmoideal could be constructed, it would be able to convey an infinite amount of information at each time step. This is in contrast to the maximum Shannon information capacity of one bit associated with each binary element.

The signum does have some advantages over the sigmoide in that it is easier to implement in hardware and much simpler to compute on a computador digital. Furthermore, the outputs of signums are binary signals which can be efficiently manipulated by computadores digitales. In a red signum with binary inputs, for instance, the output of each combinator lineal can be computed without performing weight multiplications. This involves simply adding together the values of weights with +1 inputs and




Figure 24. Example superficie ECM of signum error.

subtracting from this the values of all weights that are connected to -1 inputs.

A veces se usa un signum en una Adaline para producir decisiones de salida decisivas. The probabilidad de error is then proportional to the mean square of the error de salida $\tilde{\epsilon}$. To minimize this probabilidad de error approximately, one can easily minimize $E[(\epsilon)^2]$ instead of directly minimizing $E[(\tilde{\epsilon})^2]$ [58]. However, with only a little more computation one could minimize $E[(\tilde{\epsilon})^2]$ and typically come much closer to the objective of minimizing $E[(\tilde{\epsilon})^2]$. The sigmoide can therefore be used in training the weights even when the signum is used to form the Adaline output, as in Fig. 21.

7 Steepest-Descent Rules—Multi-Element Networks

We now study rules for descenso más pronunciado minimization of the MSE associated with entire networks of Adaline sigmoideal elements. Like their mono-elemento counterparts, the most practical and efficient descenso más pronunciado rules for redes multi-elemento typically work with one pattern presentation at a time. We will describe two descenso más pronunciado rules for multi-element redes sigmoideales, retropropagación and Regla Madaline III.

7.1 Retropropagación for Networks

The publication of the técnica de retropropagación by Rumelhart *et al.* [42] has unquestionably been the most influential development in the field of neural networks during the past decade. In retrospect, the technique seems simple. Nonetheless, largely because early neural network research dealt almost exclusively with de límite estricto nonlinearities, the idea never occurred to neural network researchers throughout the 1960s.

Los conceptos básicos de la retropropagación se comprenden fácilmente. Unfortunately, these simple ideas are often obscured by relatively intricate notation, so formal derivations of the regla de retropropagación are often tedious. We present an informal derivation of the algorithm and illustrate how it works for the simple network shown in Fig. 25.

The técnica de retropropagación is a nontrivial generalización of the single sigmoide Adaline case of Section VI-B. When ap-




Figure 25. Example two-layer red de retropropagación architecture.

plied to redes multi-elemento, the técnica de retropropagación adjusts the weights in the direction opposite the instantaneous error gradient:

$$\hat{\nabla}_k = \frac{\partial \epsilon_k^2}{\partial \mathbf{W}_k} = \begin{bmatrix} \frac{\partial \epsilon_k^2}{\partial w_{1k}} \\ \vdots \\ \frac{\partial \epsilon_k^2}{\partial w_{mk}} \end{bmatrix}. \quad (68)$$

Now, however, $\mathbf{W}_k$ is a long m -component vector of all weights in the entire network. The instantaneous error cuadrático suma ϵ_k^2 is the sum of the squares of the errors at each of the N_y outputs of the network. Thus

$$\epsilon_k^2 = \sum_{j=1}^{N_y} \epsilon_{jk}^2. \quad (69)$$

In the network example shown in Fig. 25, the error cuadrático suma is given by

$$\epsilon^2 = (d_1 - y_1)^2 + (d_2 - y_2)^2 \quad (70)$$

where we now suppress the time index k for convenience.

In its simplest form, entrenamiento por retropropagación begins by presenting an input vector de patrón X to the network, sweeping forward through the system to generate an respuesta de salida vector Y , and computing the errors at each output. The

next step involves sweeping the effects of the errors backward through the network to associate a “error cuadrático derivativo” δ with each Adaline, computing a gradient from each δ , and finally updating the weights of each Adaline based upon the corresponding gradient. A new pattern is then presented and the process is repeated. Los valores iniciales de los pesos normalmente se establecen en números aleatorios pequeños. The algorithm will not work properly with redes multicapa if the pesos iniciales are either zero or poorly chosen nonzero values.²⁰

We can get some idea about what is involved in the calculations associated with the algoritmo de retropropagación by examining the network of Fig. 25. Each of the five large circles represents a combinador lineal, as well as some associated signal paths for error retropropagación, and the corresponding adaptive machinery for updating the weights. This detail is shown in Fig. 26. The solid lines in these diagrams represent forward signal paths through the network, and the dotted lines represent the separate backward paths that are used in association with calculations of the error cuadrático derivatives δ . From Fig. 25, we see that the calculations associated with the barrido hacia atrás are of a complexity roughly equal to that represented by the

²⁰Recently, Nguyen has discovered that a more sophisticated choice of peso inicial values in capas ocultas can lead to reduced problems with local optima and dramatic increases in network training speed [100]. Experimental evidence suggests that it is advisable to choose the pesos iniciales of each capa oculta in a quasi-random manner, which ensures that at each position in a layer's input space the outputs of all but a few of its Adalines will be saturated, while ensuring that each Adaline in the layer is unsaturated in some region of its input space. When this method is used, the weights in the output layer are set to small random values.




Fig. 26. Detail of linear combiner and associated circuitry in backpropagation network.

Figure 26. Detail of combinator lineal and associated circuitry in retro-propagación network.

pasada hacia adelante through the network. The barrido hacia atrás requires the same number of function calculations as the forward sweep, but no weight multiplications in the first layer.

As stated earlier, after a pattern has been presented to the network, and the response error of each output has been calculated, the next step of the algoritmo de retropropagación involves finding the instantaneous square-error derivative δ associated with each punto de suma in the network. The error cuadrático derivativa associated with the j th Adaline in layer l is defined as²¹

$$\delta_j^{(l)} \triangleq -\frac{1}{2} \frac{\partial \epsilon^2}{\partial s_j^{(l)}}. \quad (71)$$

Each of these derivatives in essence tells us how sensitive the sum square output error of the network is to changes in the salida lineal of the associated Adaline element.

The instantaneous square-error derivatives are first computed for each element in the capa de salida. The calculation is simple. As an example, below we derive the required expression for $\delta_1^{(2)}$, the derivative associated with the top Adaline element in the capa de salida of Fig. 25. We begin with the definition of $\delta_1^{(2)}$ from Eq. (71)

$$\delta_1^{(2)} \triangleq -\frac{1}{2} \frac{\partial \epsilon^2}{\partial s_1^{(2)}}. \quad (72)$$

Expanding the squared-error term ϵ^2 by Eq. (70) yields

$$\delta_1^{(2)} = -\frac{1}{2} \frac{\partial ((d_1 - y_1)^2 + (d_2 - y_2)^2)}{\partial s_1^{(2)}} \quad (73)$$

$$= -\frac{1}{2} \frac{\partial (d_1 - \text{sgm}(s_1^{(2)}))^2}{\partial s_1^{(2)}} - \frac{1}{2} \frac{\partial (d_2 - \text{sgm}(s_1^{(2)}))^2}{\partial s_1^{(2)}}. \quad (74)$$

²¹In Fig. 25, all notation follows the convention that superscripts within parentheses indicate the layer number of the associated Adaline or input node, while subscripts identify the associated Adaline(s) within a layer.

We note that the second term is zero. Accordingly,

$$\delta_1^{(2)} = -\frac{1}{2} \frac{\partial (d_1 - \text{sgm}(s_1^{(2)}))^2}{\partial s_1^{(2)}}. \quad (75)$$

Observing that d_1 and $s_1^{(2)}$ are independent yields

$$\delta_1^{(2)} = -(d_1 - \text{sgm}(s_1^{(2)})) \frac{\partial (-\text{sgm}(s_1^{(2)}))}{\partial s_1^{(2)}} \quad (76)$$

$$= (d_1 - \text{sgm}(s_1^{(2)})) \text{sgm}'(s_1^{(2)}). \quad (77)$$

We denote the error $d_1 - \text{sgm}(s_1^{(2)})$, by $\epsilon_1^{(2)}$. Therefore,

$$\delta_1^{(2)} = \epsilon_1^{(2)} \text{sgm}'(s_1^{(2)}). \quad (78)$$

Note that this corresponds to the computation of $\delta_1^{(2)}$ as illustrated in Fig. 25. The value of δ associated with the other output element in the figure can be expressed in an analogous fashion. Thus each square-error derivative δ in the capa de salida is computed by multiplying the error de salida associated with that element by the derivative of the associated sigmoideal nonlinearity. Note from Eq. (55) that if the función sigmoideal is the hyperbolic tangent, Eq. (78) becomes simply

$$\delta_1^{(2)} = \epsilon_1^{(2)} (1 - (y_1)^2). \quad (79)$$

Developing expressions for the square-error derivatives associated with hidden layers is not much more difficult (refer to Fig. 25). We need an expression for $\delta_1^{(1)}$, the square-error derivative associated with the top element in the first layer of Fig. 25. The derivative $\delta_1^{(1)}$ is defined by

$$\delta_1^{(1)} \triangleq -\frac{1}{2} \frac{\partial \epsilon^2}{\partial s_1^{(1)}}. \quad (80)$$

Expanding this by the regla de la cadena, noting that ϵ^2 is determined entirely by the values of $s_1^{(2)}$ and $s_2^{(2)}$, yields

$$\delta_1^{(1)} = -\frac{1}{2} \left(\frac{\partial \epsilon^2}{\partial s_1^{(2)}} \frac{\partial s_1^{(2)}}{\partial s_1^{(1)}} + \frac{\partial \epsilon^2}{\partial s_2^{(2)}} \frac{\partial s_2^{(2)}}{\partial s_1^{(1)}} \right). \quad (81)$$

Using the definitions of $\delta_1^{(2)}$ and $\delta_2^{(2)}$, and then substituting expanded versions of Adaline salida lineals $s_1^{(2)}$ and $s_2^{(2)}$ gives

$$\delta_1^{(1)} = \delta_1^{(2)} \frac{\partial s_1^{(2)}}{\partial s_1^{(1)}} + \delta_2^{(2)} \frac{\partial s_2^{(2)}}{\partial s_1^{(1)}} \quad (82)$$

$$= \delta_1^{(2)} \frac{\partial}{\partial s_1^{(1)}} \left(w_{10}^{(2)} + \sum_{i=1}^3 w_{1i}^{(2)} \text{sgm}(s_i^{(1)}) \right) + \delta_2^{(2)} \frac{\partial}{\partial s_1^{(1)}} \left(w_{20}^{(2)} + \sum_{i=1}^3 w_{2i}^{(2)} \text{sgm}(s_i^{(1)}) \right). \quad (83)$$

Noting that $\partial [\text{sgm}(s_i^{(1)})] / \partial s_1^{(1)} = 0$, $i \neq j$, leaves

$$\delta_1^{(1)} = \delta_1^{(2)} w_{11}^{(2)} \text{sgm}'(s_1^{(1)}) + \delta_2^{(2)} w_{21}^{(2)} \text{sgm}'(s_1^{(1)}) \quad (84)$$

$$= [\delta_1^{(2)} w_{11}^{(2)} + \delta_2^{(2)} w_{21}^{(2)}] \text{sgm}'(s_1^{(1)}). \quad (85)$$

Now, we make the following definition:

$$\epsilon_1^{(1)} \triangleq \delta_1^{(2)} w_{11}^{(2)} + \delta_2^{(2)} w_{21}^{(2)}. \quad (86)$$

Accordingly,

$$\delta_1^{(1)} = \epsilon_1^{(1)} \text{sgm}'(s_1^{(1)}). \quad (87)$$

Referring to Fig. 25, we can trace through the circuit to verify that $\delta_1^{(1)}$ is computed in accord with Eqs. (86) and (87).

The easiest way to find values of δ for all the Adaline elements in the network is to follow the schematic diagram of Fig. 25.

Thus, the procedure for finding $\delta^{(l)}$, the square-error derivative associated with a given Adaline in capa oculta l , involves respectively multiplying each derivative $\delta^{(l+1)}$ associated with each element in the layer immediately downstream from a given Adaline by the weight that connects it to the given Adaline. These weighted square-error derivatives are then added together, producing an error term $\epsilon^{(l)}$, which, in turn, is multiplied by $\text{sgm}'(s^{(l)})$, the derivative of the given Adaline's función sigmoidee at its current punto de operación. If a network has more than two layers, this process of backpropagating the instantaneous square-error derivatives from one layer to the immediately preceding layer is successively repeated until a square-error derivative δ is computed for each Adaline in the network. This is easily shown at each layer by repeating the regla de la cadena argument associated with Eq. (81).

We now have a general method for finding a derivative δ for each Adaline element in the network. The next step is to use these δ 's to obtain the corresponding gradients. Consider an Adaline somewhere in the network which, during presentation k , has a vector de pesos $\mathbf{W}_k$, an vector de entrada $\mathbf{X}_k$, and a linear output $s_k = \mathbf{W}_k^T \mathbf{X}_k$. The gradiente instantáneo for this Adaline element is

$$\hat{\nabla}_k = \frac{\partial \epsilon_k^2}{\partial \mathbf{W}_k}. \quad (88)$$

This can be written as

$$\hat{\nabla}_k = \frac{\partial \epsilon_k^2}{\partial \mathbf{W}_k} = \frac{\partial \epsilon_k^2}{\partial s_k} \frac{\partial s_k}{\partial \mathbf{W}_k}. \quad (89)$$

Note that $\mathbf{W}_k$ and $\mathbf{X}_k$ are independent so

$$\frac{\partial s_k}{\partial \mathbf{W}_k} = \frac{\partial \mathbf{W}_k^T \mathbf{X}_k}{\partial \mathbf{W}_k} = \mathbf{X}_k. \quad (90)$$

Therefore,

$$\hat{\nabla}_k = \frac{\partial \epsilon_k^2}{\partial s_k} \mathbf{X}_k. \quad (91)$$

For this element,

$$\delta_k = -\frac{1}{2} \frac{\partial \epsilon_k^2}{\partial s_k}. \quad (92)$$

Accordingly,

$$\hat{\nabla}_k = -2\delta_k \mathbf{X}_k. \quad (93)$$

Updating the weights of the Adaline element using the method of descenso más pronunciado with the gradiente instantáneo is a process represented by

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \mu(-\hat{\nabla}_k) = \mathbf{W}_k + 2\mu\delta_k \mathbf{X}_k. \quad (94)$$

Thus, after backpropagating all square-error derivatives, we complete a retropropagación iteration by adding to each vector de pesos the corresponding input vector scaled by the associated square-error derivative. Eq. (94) and the means for finding δ_k comprise the general actualización de pesos rule of the retropropagación algorithm.

There is a great similarity between Eq. (94) and the μ -algorithm LMS (33), but one should view this similarity with caution. The quantity δ_k , defined as a squared-error derivative, might appear to play the same role in retropropagación as that played by the error in the μ -algorithm LMS. However, δ_k is not an error. Adaptation of the given Adaline is effected to reduce the squared output error ϵ_k^2 , not δ_k of the given Adaline or of any other Adaline in the network. The objective is not to reduce the δ_k 's of the network, but to reduce ϵ_k^2 at the network output.

It is interesting to examine the actualización de pesos that retropropagación imposes on the Adaline elements in the capa de salida. Substituting Eq. (77) into Eq. (94) reveals the Adaline which provides output y_1 in Fig. 25 is updated by the rule

$$\mathbf{W}_{k+1} = \mathbf{W}_k + 2\mu\epsilon_1^{(2)} \text{sgm}'(s_1^{(2)}) \mathbf{X}_k. \quad (95)$$

This rule turns out to be identical to the single Adaline version (54) of the regla de retropropagación. This is not surprising since the output Adaline is provided with both input signals and respuestas deseadas, so its training circumstance is the same as that experienced by an Adaline trained in isolation.

There are many variants of the algoritmo de retropropagación. Sometimes, the size of μ is reduced during training to diminish the effects of ruido de gradiente in the weights. Another extension is the técnica de momento [42] which involves including in the cambio de pesos vector $\Delta \mathbf{W}_k$ of each Adaline a term proportional to the corresponding cambio de pesos from the previous iteration. That is, Eq. (94) is replaced by a pair of equations:

$$\Delta \mathbf{W}_k = 2\mu(1 - \eta)\delta_k \mathbf{X}_k + \eta\Delta \mathbf{W}_{k-1} \quad (96)$$

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \Delta \mathbf{W}_k. \quad (97)$$

where the constante de momento $0 \leq \eta < 1$ is in practice usually set to something around 0.8 or 0.9.

The técnica de momento filtros pasa-bajos the actualización de pesos and thereby tends to resist erratic cambios de pesos caused either by ruido de gradiente or high spatial frequencies in the superficie ECM. The factor $(1 - \eta)$ in Eq. (96) is included to give the filter a ganancia DC of unity so that the tasa de aprendizaje μ does not need to be stepped down as the constante de momento η is increased. A momentum term can also be added to the update equations of other algorithms discussed in this paper. A detailed analysis of stability issues associated with momentum updating for the μ -algorithm LMS, for instance, has been described by Shynk and Roy [124].

In our experience, the técnica de momento used alone is usually of little value. We have found, however, that it is often useful to apply the technique in situations that require relatively "clean"²² estimación del gradientes. One case is a normalized actualización de pesos equation which makes the network's vector de pesos move the same distancia euclidiana with each iteration. This can be accomplished by replacing Eq. (96) and (97)

²²"Clean" estimación del gradientes are those with little ruido de gradiente.

with

$$\Delta_k = \delta_k \mathbf{X}_k + \eta \Delta_{k+1} \quad (98)$$

$$\mathbf{W}_{k+1} = \mathbf{W}_k + \frac{\mu \Delta_k}{\sqrt{\sum_{\text{all Adalines}} |\Delta_k|^2}}. \quad (99)$$

where again $0 < \eta < 1$. The actualización de pesos determined by Eqs. (98) and (99) can help a network find a solution when a relatively flat local region in the MSE surface is encountered. The weights move by the same amount whether the surface is flat or inclined. It is reminiscent of α -LMS because the gradient term in the actualización de pesos equation is normalized by a time-varying factor. The actualización de pesos rule could be further modified by including terms from both techniques associated with Eqs. (96) through (99). Other methods for speeding up entrenamiento por retropropagación include Fahlman's popular quickprop method [125], as well as the delta-bar-delta approach reported in an excellent paper by Jacobs [126].²³

One of the most promising new areas of neural network research involves retropropagación variants for training various recurrent (retroalimentación de señal) networks. Recently, regla de retropropagacións have been derived for training recurrent networks to learn static associations [127, 128]. More interesting is the on-line technique of Williams and Zipser [129] which allows a wide class of redes recurrentes to learn dynamic associations and trajectories. A more general and computationally viable variant of this technique has been advanced by Narendra and Parthasarathy [104]. These on-line methods are generalizaciones of a well-known descenso más pronunciado algorithm for training linear IIR filters [30, 130].

An equivalent technique that is usually far less computationally intensive but best suited for off-line computation [37, 42, 131], called "retropropagación through time," has been used by Nguyen and Widrow [50] to enable a neural network to learn without a teacher how to back up a computer-simulated camión con remolque to a plataforma de carga (Fig. 27). This is a highly dirección no lineal task and it is not yet known how to design a controller to perform it. Nevertheless, with just 6 inputs providing information about the current position of the truck, a two-layer neural network with only 26 Adalines was able to learn of its own accord to solve this problem. Once trained, the network could successfully back up the truck from any initial position and orientation in front of the plataforma de carga.

7.2 Regla Madaline III for Networks

It is difficult to build neural networks with hardware analógico that can be trained effectively by the popular técnica de retropropagación. Attempts to overcome this difficulty have led to

²³Jacob's paper, like many other papers in the literature, assumes for analysis that the gradiente verdaderos rather than gradiente instantáneos are used to update the weights, that is, that weights are changed periodically, only after all patrones de entrenamiento are presented. This eliminates ruido de gradiente but can slow down training enormously if the conjunto de entrenamiento is large. The delta-bar-delta procedure in Jacob's paper involves monitoring changes of the gradiente verdaderos in response to cambios de pesos. It should be possible to avoid the expense of computing the true gradients explicitly in this case by instead monitoring changes in the outputs of, say, two momentum filters with different time constants.




Figure 27. Example truck backup sequence.

the development of the MRIII algorithm. A commercial analog neurocomputing chip based primarily on this algorithm has already been devised [132]. The method described in this section is a generalización of the single Adaline MRIII technique (60). The multi-element generalización of the other single element MRIII rule (61) is described in [133].

The MRIII algorithm can be readily described by referring to Fig. 28. Although this figure shows a simple two-layer feedforward architecture, the procedure to be developed will work for neural networks with any number of Adaline elements in any feedforward structure. In [133], we discuss variants of the basic MRIII approach that allow descenso más pronunciado training to be applied to more general network topologies, even those with retroalimentación de señal.

Assume that an patrón de entrada $\mathbf{X}$ and its associated desired respuestas de salida d_1 and d_2 are presented to the network of Fig. 28. At this point, we measure the sum squared respuesta de salida error $\epsilon^2 = (d_1 - y_1)^2 + (d_2 - y_2)^2 = \epsilon_1^2 + \epsilon_2^2$. We then add a small quantity Δs to a selected Adaline in the network, providing a perturbation to the element's linear sum. This perturbation propagates through the network, and causes a change in the sum of the squares of the errors, $\Delta(\epsilon^2) = \Delta(\epsilon_1^2 + \epsilon_2^2)$. An easily measured ratio is

$$\frac{\Delta(\epsilon^2)}{\Delta s} = \frac{\Delta(\epsilon_1^2 + \epsilon_2^2)}{\Delta s} \simeq \frac{\partial(\epsilon^2)}{\partial s}. \quad (100)$$

Below we use this to obtain the gradiente instantáneo of ϵ_k^2 with respect to the vector de pesos of the selected Adaline. For the k th presentation, the gradiente instantáneo is

$$\hat{\nabla}_k = \frac{\partial(\epsilon_k^2)}{\partial \mathbf{W}_k} = \frac{\partial(\epsilon_k^2)}{\partial s_k} \frac{\partial s_k}{\partial \mathbf{W}_k} = \frac{\partial(\epsilon_k^2)}{\partial s_k} \mathbf{X}_k. \quad (101)$$




Figure 28. Example two-layer Madaline III architecture.

Replacing the derivative with a ratio of differences yields

$$\hat{\nabla}_k \simeq \frac{\Delta(\epsilon_k^2)}{\Delta s} \mathbf{X}_k. \quad (102)$$

The idea of obtaining a derivative by perturbing the salida lineal of the selected Adaline element is the same as that expressed for the single element in Section VI-B, except that here the error is obtained from the output of a red multi-elemento rather than from the output of a single element.

The gradient (102) can be used to optimize the vector de pesos in accord with the method of descenso más pronunciado:

$$\mathbf{W}_{k+1} = \mathbf{W}_k - \mu \frac{\Delta(\epsilon_k^2)}{\Delta s} \mathbf{X}_k. \quad (103)$$

Maintaining the same patrón de entrada, one could either perturb all the elements in the network in sequence, adapting after each gradient calculation, or else the derivatives could be computed and stored to allow all Adalines to be adapted at once. These two MRIII approaches both involve the same actualización de pesos equation (103), and if μ is small, both lead to equivalent solutions. With large μ , experience indicates that adapting one element at a time results in convergence after fewer iterations, especially in large networks. Storing the gradients, however, has the advantage that after the initial unperturbed error is measured during a given training presentation, each estimación del gradiente requires only the perturbed error measurement. If adaptations take place after each error measurement, both perturbed and unperturbed errors must be measured for each gradient calculation. This is because each actualización de pesos changes the associated unperturbed error.

7.3 Comparison of MRIII with MRII

MRIII se derivó de MRII reemplazando las no linealidades signum por sigmoideas. The similarity of these algorithms becomes evident when comparing Fig. 28, representing MRIII, with Fig. 16, representing MRII.

La red MRII es altamente discontinua y no lineal. Using an instantaneous gradient to adjust the weights is not possible. In fact, from the superficie ECM for the Adaline signum presented

in Section VI-B, it is clear that even descenso por gradiente techniques that use the gradiente verdadero could run into severe problems with local minima. The idea of adding a perturbation to the linear sum of a selected Adaline element is workable, however. If the error de Hamming has been reduced by the perturbation, the Adaline is adapted to reverse its decisión de salida. This weight change is in the LMS direction, along its $\mathbf{X}$ -vector. If adapting the Adaline would not reduce network error de salida, it is not adapted. This is in accord with the principio de mínima perturbación. The Adalines selected for possible adaptation are those whose analog sums are closest to zero, that is, the Adalines that can be adapted to give opposite responses with the smallest cambios de pesos. It is useful to note that with binary ± 1 desired responses, the Hamming error is equal to $1/4$ the error cuadrático suma. Minimizing the error de salida de Hamming is therefore equivalent to minimizing the error de salida cuadrático suma.

El algoritmo MRIII funciona de manera similar. All the Adalines in the MRIII network are adapted, but those whose analog sums are closest to zero will usually be adapted most strongly, because the sigmoide has its maximum slope at zero, contributing to high gradient values. As with MRII, the objective is to change the weights for the given input presentation to reduce the error cuadrático suma at the network output. In accord with the principio de mínima perturbación, the vectores de pesos of the Adaline elements are adapted in the LMS direction, along their $\mathbf{X}$ -vectors, and are adapted in proportion to their capabilities for reducing the error cuadrático suma (the square of the error euclidiano) at the output.

7.4 Comparison of MRIII with Retropropagación

In Section VI-B, we argued that for the Adaline sigmoideal element, the MRIII algorithm (61) is essentially equivalent to the algoritmo de retropropagación (54). The same argument can be extended to the network of Adaline elements, demonstrating that if Δs is small and adaptation is applied to all elements in the network at once, then MRIII is essentially equivalent to retropropagación. That is, to the extent that the sample derivative $\Delta(\epsilon_k^2)/\Delta s$ from Eq. (103) is equal to the analytical derivative $\partial \epsilon_k^2 / \partial s_k$ from Eq. (91), the two rules follow identical gradient instantáneos, and thus perform identical actualización de pesos.

The algoritmo de retropropagación requires fewer operations than MRIII to calculate gradients, since it is able to take advantage of *a priori* knowledge of the sigmoide nonlinearities and their derivative functions. Conversely, the MRIII algorithm uses no prior knowledge about the characteristics of the sigmoide functions. Rather, it acquires gradiente instantáneos from perturbation measurements. Using MRIII, tolerances on the sigmoide implementations can be greatly relaxed compared to acceptable tolerances for successful retropropagación.

Steepest-descent training of redes multicapa implemented by computer simulation or by precise parallel digital hardware is usually best carried out by retropropagación. During each training presentation, the retropropagación method requires only one forward computation through the network followed by one backward computation in order to adapt all the weights of an entire network. To accomplish the same effect with the form of MRIII

that updates all weights at once, one measures the unperturbed error followed by a number of perturbed error measurements equal to the number of elements in the network. This could require a lot of computation.

If a network is to be implemented in hardware analógico, however, experience has shown that MRIII offers strong advantages over retropropagación. Comparison of Fig. 25 with Fig. 28 demonstrates the relative simplicity of MRIII. All the apparatus for backward propagation of error-related signals is eliminated, and the weights do not need to carry signals in both directions (see Fig. 26). MRIII is a much simpler algorithm to build and to understand, and in principle it produces the same gradiente instantáneo as the retropropagación algorithm. The técnica de momento and most other common variants of the algoritmo de retropropagación can be applied to MRIII training.

7.5 MSE Surfaces of Neural Networks

In Section VI-B, “typical” mean-square-error surfaces of sigmoide and signum Adalines were shown, indicating that Adaline sigmoideals are much more conducive to gradient approaches than Adaline signums. The same phenomena result when Adalines are incorporated into redes multi-elemento. The superficies ECM of MRII networks are reasonably chaotic and will not be explored here. In this section we examine only superficies ECM from a typical entrenamiento por retropropagación problem with a sigmoideal neural network.

In a network with more than two weights, the superficie ECM is high-dimensional and difficult to visualize. It is possible, however, to look at slices of this surface by plotting the superficie ECM created by varying two of the weights while holding all others constant. The surfaces plotted in Figs. 29 and 30 show two such slices of the superficie ECM from a typical learning problem involving, respectively, an untrained sigmoideal network and a trained one. The first surface resulted from varying two de la primera capa weights of an untrained network. The second surface resulted from varying the same two weights after the network was fully trained. The two surfaces are similar, but the second one has a deeper minimum which was carved out by the retropropagación learning process. Figs. 31 and 32 resulted from varying a different set of two weights in the same network. Fig. 31 is the result from varying a de la primera capa weight and de la tercera capa weight in the untrained network, whereas Fig. 32 is the surface that resulted from varying the same two weights after the network was trained.

By studying many plots, it becomes clear that retropropagación and MRIII will be subject to convergence on óptimos locales. Lo mismo es cierto para MRII. The most common remedy for this is the sporadic addition of noise to the weights or gradients. Some of the “recocido simulado” methods [47] do this. Another method involves retraining the network several times using different random peso inicial values until a satisfactorio solution is found.

Solutions found by people in everyday life are usually not optimal, but many of them are useful. If a local optimum yields satisfactory performance, often there is simply no need to search for a better solution.




Figure 29. Example superficie ECM of untrained sigmoideal network as a function of two de la primera capa weights.




Figure 30. Example superficie ECM of trained sigmoideal network as a function of two de la primera capa weights.

8 Summary

This year is the 30th anniversary of the publication of the regla del Perceptrón by Rosenblatt and the algoritmo LMS by Widrow and Hoff. It has also been 16 years since Werbos first published the algoritmo de retropropagación. These reglas de aprendizaje and several others have been studied and compared. Although they differ significantly from each other, they all belong to the same “family.”

Se estableció una distinción entre reglas de corrección de error y reglas de descenso más pronunciado. The former includes the regla del Perceptrón, Mays’ rules, the α -algoritmo LMS, the original Madaline I rule of 1962, and the Madaline II rule. The latter includes the μ -algoritmo LMS, the Madaline III rule, and the algoritmo de retropropagación. Fig. 33 categorizes the reglas de aprendizaje that have been studied.

Although these algorithms have been presented as established reglas de aprendizaje, one should not gain the impression that they are perfect and frozen for all time. Son posibles variaciones para cada uno de ellos. They should be regarded as substrates upon which to build new and better rules. There is a tremendous amount of invention waiting “in the wings.” Esperamos con interés los próximos 30 años.




Figure 31. Example superficie ECM of untrained sigmoideal network as a function of a de la primera capa weight and a de la tercera capa weight.

weight.




Figure 32. Example superficie ECM of trained sigmoideal network as a function of a de la primera capa weight and a de la tercera capa weight.

References

- [1] K. Steinbuch and V. A. W. Piske, "Learning matrices and their applications," *IEEE Trans. Electron. Comput.*, vol. EC-12, pp. 846–862, Dec. 1963.
- [2] B. Widrow, "Generalization and information storage in networks of adaline 'neurons'," in *Self-Organizing Systems 1962*, M. Yovitz, G. Jacobi, and G. Goldstein, Eds. Washington, DC: Spartan Books, 1962, pp. 435–461.
- [3] L. Stark, M. Okajima, and G. H. Whipple, "Computer reconocimiento de patrones techniques: Electrocardiographic diagnosis," *Commun. Ass. Comput. Mach.*, vol. 5, pp. 527–532, Oct. 1962.
- [4] F. Rosenblatt, "Two theorems of statistical separability in the perceptron," in *Mechanization of Thought Processes: Proc. Symp. National Physical Laboratory, Nov. 1958*, vol. 1, pp. 421–456. London: HM Stationery Office, 1959.
- [5] F. Rosenblatt, *Principles of Neurodynamics: Perceptrons and the Theory of Brain Mechanisms*. Washington, DC: Spartan Books, 1962.




Fig. 33. Learning rules.

Figure 33. Learning rules.

- [6] C. von der Malsburg, "Self-organizing of orientation sensitive cells in the striate cortex," *Kybernetik*, vol. 14, pp. 85–100, 1973.
- [7] S. Grossberg, "Adaptive pattern classification and universal re-coding, I: Parallel development and coding of neural detectors de características," *Biolog. Cybernetics*, vol. 23, pp. 121–134, 1976.
- [8] K. Fukushima, "Cognitrón: A autoorganizado multilayered neural network," *Biolog. Cybernetics*, vol. 20, pp. 121–136, 1975.
- [9] K. Fukushima, "Neocognitrón: A autoorganizado neural network model for a mechanism of reconocimiento de patrones unaffected by shift in position," *Biolog. Cybernetics*, vol. 36, pp. 193–202, 1980.
- [10] B. Widrow, "Bootstrap learning in lógica de umbral systems," presented at the American Automatic Control Council (Theory Committee), IFAC Meeting, London, England, June 1966.
- [11] B. Widrow, N. K. Gupta, and S. Maitra, "Punish/reward: Learning with a critic in adaptive threshold systems," *IEEE Trans. Syst., Man, Cybernetics*, vol. SMC-3, pp. 455–465, Sept. 1973.
- [12] A. G. Barto, R. S. Sutton, and C. W. Anderson, "Neuronlike adaptive elements that can solve difficult learning control problems," *IEEE Trans. Syst., Man, Cybernetics*, vol. SMC-13, pp. 834–846, 1983.
- [13] J. S. Albus, "A new approach to manipulator control: the cerebellar model articulation controller (CMAC)," *J. Dyn. Sys., Meas., Contr.*, vol. 97, pp. 220–227, 1975.
- [14] W. T. Miller, III, "Sensor-based control of robotic manipulators using a general algoritmo de aprendizaje," *IEEE J. Robotics Automat.*, vol. RA-3, pp. 157–165, Apr. 1987.
- [15] S. Grossberg, "Adaptive pattern classification and universal re-coding, II: Feedback, expectation, olfaction, and illusions," *Biolog. Cybernetics*, vol. 23, pp. 187–202, 1976.
- [16] G. A. Carpenter and S. Grossberg, "A massively parallel architecture for a autoorganizado neural reconocimiento de patrones machine," *Computer Vision, Graphics, and Image Processing*, vol. 37, pp. 54–115, 1983.

- [17] G. A. Carpenter and S. Grossberg, "Art 2: Self-organization of stable category recognition codes for analog output patterns," *Applied Optics*, vol. 26, pp. 4919–4930, Dec. 1, 1987.
- [18] G. A. Carpenter and S. Grossberg, "Art 3 hierarchical search: Chemical transmitters in autoorganizado reconocimiento de patrones architectures," in *Proc. Int. Joint Conf. on Neural Networks*, vol. 2, pp. 30–33, Wash., DC, Jan. 1990.
- [19] T. Kohonen, "Self-organized formation of topologically correct mapas de características," *Biolog. Cybernetics*, vol. 43, pp. 59–69, 1982.
- [20] T. Kohonen, *Self-organization and Associative Memory*. New York: Springer-Verlag, 2d ed., 1988.
- [21] D. O. Hebb, *The Organization of Behavior*. New York: Wiley, 1949.
- [22] J. J. Hopfield, "Neural networks and physical systems with emergent collective computational abilities," *Proc. Natl. Acad. Sci.*, vol. 79, pp. 2554–2558, Apr. 1982.
- [23] J. J. Hopfield, "Neurons with graded response have collective computational properties like those of two-state neurons," *Proc. Natl. Acad. Sci.*, vol. 81, pp. 3088–3092, May 1984.
- [24] B. Kosko, "Adaptive bidirectional associative memories," *Appl. Optics*, vol. 26, pp. 4947–4960, Dec. 1, 1987.
- [25] G. E. Hinton, R. J. Sejnowski, and D. H. Ackley, "Boltzmann machines: Constraint satisfaction networks that learn," Tech. Rep. CMU-CS-84-119, Carnegie-Mellon University, Dept. of Computer Science, 1984.
- [26] G. E. Hinton and T. J. Sejnowski, "Learning and relearning in Boltzmann machines," in *Parallel Distributed Processing*, vol. 1, ch. 7, D. E. Rumelhart and J. L. McClelland, Eds. Cambridge, MA: M.I.T. Press, 1986.
- [27] L. R. Talbert *et al.*, "A real-time adaptive speech-recognition system," Tech. rep., Stanford University, 1963.
- [28] M. J. C. Hu, *Application of the Adaline System to Weather Forecasting*. Thesis, Tech. Rep. 6775-1, Stanford Electron. Labs., Stanford, CA, June 1964.
- [29] B. Widrow, "The original adaptive neural net broom-balancer," *Proc. IEEE Intl. Symp. Circuits and Systems*, pp. 351–357, Phila., PA, May 4–7 1987.
- [30] B. Widrow and S. D. Stearns, *Adaptive Signal Processing*. Englewood Cliffs, NJ: Prentice-Hall, 1985.
- [31] B. Widrow, P. Mantey, L. Griffiths, and B. Goode, "Adaptive antenna systems," *Proc. IEEE*, vol. 55, pp. 2143–2159, Dec. 1967.
- [32] B. Widrow, "Adaptive inverse control," *Proc. 2d Intl. Fed. of Automatic Control Workshop*, pp. 1–5, Lund, Sweden, July 1–3, 1986.
- [33] B. Widrow *et al.*, "Adaptive noise cancelling: Principles and applications," *Proc. IEEE*, vol. 63, pp. 1692–1716, Dec. 1975.
- [34] R. W. Lucky, "Automatic equalization for digital communication," *Bell Syst. Tech. J.*, vol. 44, pp. 547–588, Apr. 1965.
- [35] R. W. Lucky *et al.*, *Principles of Data Communication*. New York: McGraw-Hill, 1968.
- [36] M. M. Sondhi, "An cancelador de eco adaptativo," *Bell Syst. Tech. J.*, vol. 46, pp. 497–511, Mar. 1967.
- [37] P. Werbos, *Beyond Regression: New Tools for Prediction and Analysis in the Behavioral Sciences*. tesis doctoral, Harvard University, Cambridge, MA, Aug. 1974.
- [38] Y. le Cun, "A theoretical framework for back-propagation," in *Proc. 1988 Connectionist Models Summer School*, D. Touretzky, G. Hinton, and T. Sejnowski, Eds. June 17–26, pp. 21–28. San Mateo, CA: Morgan Kaufmann.
- [39] D. Parker, "Learning-logic," Invention Report S81-64, File 1, Office of Technology Licensing, Stanford University, Stanford, CA, Oct. 1982.
- [40] D. Parker, "Learning-logic," Technical Report TR-47, Center for Computational Research in Economics and Management Science, M.I.T., Apr. 1985.
- [41] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning representaciones internas by error propagation," ICS Report 8506, Institute for Cognitive Science, UCSD, La Jolla, CA, Sept. 1985.
- [42] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning representaciones internas by error propagation," in *Parallel Distributed Processing*, vol. 1, ch. 8, D. E. Rumelhart and J. L. McClelland, Eds. Cambridge, MA: M.I.T. Press, 1986.
- [43] B. Widrow, R. G. Winter, and R. Baxter, "Learning phenomena in layered neural networks," *Proc. 1st IEEE Intl. Conf. on Neural Networks*, vol. 2, pp. 411–429, San Diego, CA, June 1987.
- [44] R. P. Lippmann, "An introduction to computing with neural nets," *IEEE ASSP Mag.*, Apr. 1987.
- [45] J. A. Anderson and E. Rosenfeld, Eds., *Neurocomputing: Foundations of Research*. Cambridge, MA: M.I.T. Press, 1988.
- [46] N. Nilsson, *Learning Machines*. New York: McGraw-Hill, 1965.
- [47] D. E. Rumelhart and J. L. McClelland, Eds., *Parallel Distributed Processing*. Cambridge, MA: M.I.T. Press, 1986.
- [48] B. Moore, "Art 1 and agrupamiento de patrones," in *Proc. 1988 Connectionist Models Summer School*, D. Touretzky, G. Hinton, and T. Sejnowski, Eds., June 17–26 1988, pp. 174–185, San Mateo, CA: Morgan Kaufmann.
- [49] *DARPA Neural Network Study*. Fairfax, VA: AFCEA International Press, 1988.
- [50] D. Nguyen and B. Widrow, "The truck backer-upper: An example of self-learning in neural networks," *Proc. Intl. Joint Conf. on Neural Networks*, vol. 2, pp. 357–363, Wash., DC, June 1989.
- [51] T. J. Sejnowski and C. R. Rosenberg, "Nettalk: a parallel network that learns to read aloud," Tech. Rep. JHU/EECS-86/01, Johns Hopkins University, 1986.
- [52] T. J. Sejnowski and C. R. Rosenberg, "Parallel networks that learn to pronounce English text," *Complex Systems*, vol. 1, pp. 145–168, 1987.

- [53] P. M. Shea and V. Lin, "Detection of explosives in checked airline baggage using an artificial neural system," *Proc. Intl. Joint Conf. on Neural Networks*, vol. 2, pp. 31–34, Wash., DC, June 1989.
- [54] D. G. Bounds, P. J. Lloyd, B. Mathew, and G. Waddell, "A multilayer perceptron network for the diagnosis of low back pain," *Proc. 2d IEEE Intl. Conf. on Neural Networks*, vol. 2, pp. 481–489, San Diego, CA, July 1988.
- [55] G. Bradshaw, R. Fozzard, and L. Ceci, "A connectionist expert system that actually works," in *Advances in Neural Information Processing Systems I*, D. S. Touretzky, Ed. San Mateo, CA: Morgan Kaufmann, 1989, pp. 248–255.
- [56] N. Mokhoff, "Neural nets making the leap out of lab," *Electronic Engineering Times*, p. 1, Jan. 22, 1990.
- [57] C. A. Mead, *Analog VLSI and Neural Systems*. Reading, MA: Addison-Wesley, 1989.
- [58] B. Widrow and M. E. Hoff, Jr., "Adaptive switching circuits," *1960 IRE Western Electric Show and Convention Record*, Part 4, pp. 96–104, Aug. 23, 1960.
- [59] B. Widrow and M. E. Hoff, Jr., "Adaptive switching circuits," Tech. Rep. 1553-1, Stanford Electron. Labs., Stanford, CA, June 30, 1960.
- [60] P. M. Lewis II and C. Coates, *Threshold Logic*. New York: Wiley, 1967.
- [61] T. M. Cover, *Geometrical and Statistical Properties of Linear Threshold Devices*. tesis doctoral, Tech. Rep. 6107-1, Stanford Electron. Labs., Stanford, CA, May 1964.
- [62] R. J. Brown, *Adaptive Multiple-Output Threshold Systems and Their Storage Capacities*. Thesis, Tech. Rep. 6771-1, Stanford Electron. Labs., Stanford, CA, June 1964.
- [63] R. O. Winder, *Threshold Logic*. tesis doctoral, Princeton University, Princeton, NJ, 1962.
- [64] S. H. Cameron, "An estimate of the complexity requisite in a universal decision network," *Proc. 1960 Bionics Symposium*, Wright Air Development Division Tech. Rep. 60-600, pp. 197–211, Dayton, OH, Dec. 1960.
- [65] R. D. Joseph, "The number of orthants in n -space intersected by an s -dimensional subspace," Tech. Memorandum 8, Project PARA, Cornell Aeronautical Laboratory, Buffalo, NY, 1960.
- [66] D. F. Specht, *Generation of Polynomial Discriminant Functions for Pattern Recognition*. tesis doctoral, Tech. Rep. 6764-5, Stanford Electron. Labs., Stanford, CA, May 1966.
- [67] D. F. Specht, "Vectorcardiographic diagnosis using the método discriminante polinomial de reconocimiento de patrones," *IEEE Trans. Biomed. Eng.*, vol. BME-14, pp. 90–95, Apr. 1967.
- [68] D. F. Specht, "Generation of discriminante polinomial functions for reconocimiento de patrones," *IEEE Trans. Electron. Comput.*, vol. EC-16, pp. 308–319, June 1967.
- [69] A. R. Barron, "Adaptive learning networks: Development and application in the United States of algorithms related to gmdh," in *Self-Organizing Methods in Modeling*, S. J. Farlow, Ed., New York: Marcel Dekker Inc., 1984, pp. 25–65.
- [70] A. R. Barron, "Predicted error cuadrático: A criterion for automatic model selection," in *Self-Organizing Methods in Modeling*, S. J. Farlow, Ed. New York: Marcel Dekker Inc., 1984, pp. 87–103.
- [71] A. R. Barron and R. L. Barron, "Statistical learning networks: A unifying view," *1988 Symp. on the Interface: Statistics and Computing Science*, pp. 192–203, Reston, VA, Apr. 21–23, 1988.
- [72] A. G. Ivakhnenko, "Polynomial theory of complex systems," *IEEE Trans. Syst., Man, Cybernetics*, SMC-1, pp. 364–378, Oct. 1971.
- [73] Y. H. Pao, "Functional link nets: Removing capas ocultas," *AI Expert*, pp. 60–68, Apr. 1989.
- [74] C. L. Giles and T. Maxwell, "Learning, invariance, and generalización in high-order neural networks," *Applied Optics*, vol. 26, pp. 4972–4978, Dec. 1, 1987.
- [75] M. E. Hoff, Jr., *Learning Phenomena in Networks of Adaptive Switching Circuits*. tesis doctoral, Tech. Rep. 1554-1, Stanford Electron. Labs., Stanford, CA, July 1962.
- [76] W. C. Ridgway III, *An Adaptive Logic System with Generalizing Properties*. tesis doctoral, Tech. Rep. 1556-1, Stanford Electron. Labs., Stanford, CA, April 1962.
- [77] F. H. Glanz, *Statistical Extrapolation in Certain Adaptive Pattern-Recognition Systems*. tesis doctoral, Tech. Rep. 6767-1, Stanford Electron. Labs., Stanford, CA, May 1965.
- [78] B. Widrow, "Adaline and Madaline—1963, plenary speech," *Proc. 1st IEEE Intl. Conf. on Neural Networks*, vol. 1, pp. 145–158, San Diego, CA, June 23, 1987.
- [79] B. Widrow, "An adaptive 'adaline' neuron using chemical 'memistors'," Tech. Rep. 1553-2, Stanford Electron. Labs., Stanford, CA, Oct. 17, 1960.
- [80] C. L. Giles, R. D. Griffin, and T. Maxwell, "Encoding geometric invariances in higher order neural networks," in *Neural Information Processing Systems*, D. Z. Anderson, Ed. New York: American Institute of Physics, 1988, pp. 301–309.
- [81] D. Casasent and D. Psaltis, "Position, rotation, and invariante a escala optical correlation," *Appl. Optics*, vol. 15, pp. 1795–1799, July 1976.
- [82] W. L. Reber and J. Lyman, "An artificial neural system design for rotation and invariante a escala reconocimiento de patrones," *Proc. 1st IEEE Intl. Conf. on Neural Networks*, vol. 4, pp. 277–283, San Diego, CA, June 1987.
- [83] B. Widrow and R. G. Winter, "Neural nets for filtrado adaptativo and adaptive reconocimiento de patrones," *IEEE Computer*, pp. 25–39, Mar. 1988.
- [84] A. Khotanzad and Y. H. Hong, "Rotation invariant reconocimiento de patrones using zernike moments," *Proc. 9th Intl. Conf. on Pattern Recognition*, vol. 1, pp. 326–328, 1988.
- [85] C. von der Malsburg, "Pattern recognition by labeled emparejamiento de grafos," *Neural Networks*, vol. 1, pp. 141–148, 1988.

- [86] A. Waibel, T. Hanazawa, G. Hinton, K. Shikano, and K. J. Lang, "Phoneme recognition using time delay neural networks," *IEEE Trans. Acoust., Speech, and Signal Processing*, vol. ASSP-37, pp. 328–339, Mar. 1989.
- [87] C. M. Newman, "Memory capacity in neural network models: Rigorous lower bounds," *Neural Networks*, vol. 1, pp. 223–238, 1988.
- [88] Y. S. Abu-Mostafa and J. St. Jacques, "Information capacity of the hopfield model," *IEEE Trans. Inform. Theory*, vol. IT-31, pp. 461–464, 1985.
- [89] Y. S. Abu-Mostafa, "Neural networks for computing?" in *Neural Networks for Computing*, Amer. Inst. of Phys. Conf. Proc. No. 151, J. S. Denker, Ed. New York: American Institute of Physics, 1986, pp. 1–6.
- [90] S. S. Venkatesh, "Epsilon capacity of neural networks," in *Neural Networks for Computing*, Amer. Inst. of Phys. Conf. Proc. No. 151, J. S. Denker, Ed. New York: American Institute of Physics, 1986, pp. 440–445.
- [91] J. D. Greenfield, *Practical Digital Design Using IC's*. 2d ed., New York: Wiley, 1983.
- [92] M. Stinchcombe and H. White, "Universal approximation using redes con alimentación hacia adelante with non-sigmoide capa oculta activation functions," *Proc. Intl. Joint Conf. on Neural Networks*, vol. 1, pp. 613–617, Wash., DC, June 1989.
- [93] G. Cybenko, "Continuous valued neural networks with two capas ocultas are sufficient," Tech. Rep., Dept. of Computer Science, Tufts University, Mar. 1988.
- [94] B. Irie and S. Miyake, "Capabilities of three-layered perceptrons," *Proc. 2d IEEE Intl. Conf. on Neural Networks*, vol. 1, pp. 641–647, San Diego, CA, July 1988.
- [95] M. L. Minsky and S. A. Papert, *Perceptrons: An Introduction to Computational Geometry*. Cambridge, MA: M.I.T. Press, expanded ed., 1988.
- [96] M. W. Roth, "Survey of neural network technology for automatic reconocimiento de blancos," *IEEE Trans. Neural Networks*, vol. 1, pp. 28–43, Mar. 1990.
- [97] T. M. Cover, "Capacity problems for linear machines," in *Pattern Recognition*, L. N. Kanal, Ed. Wash., DC: Thompson Book Co., 1968, pp. 283–289, part 3.
- [98] E. B. Baum, "On the capabilities of multilayer perceptrons," *J. Complexity*, vol. 4, pp. 193–215, Sept. 1988.
- [99] A. Lapedes and R. Farber, "How neural networks work," Tech. Rep. LA-UR-88-418, Los Alamos Nat. Laboratory, Los Alamos, NM, 1987.
- [100] D. Nguyen and B. Widrow, "Improving the learning speed of 2-layer neural networks by choosing initial values of the adaptive weights," *Proc. Intl. Joint Conf. on Neural Networks*, San Diego, CA, June 1990.
- [101] G. Cybenko, "Approximation by superpositions of a sigmoideal function," *Mathematics of Control, Signals, and Systems*, vol. 2, 1989.
- [102] E. B. Baum and D. Haussler, "What size net gives valid generalización?" *Neural Computation*, vol. 1, pp. 151–160, 1989.
- [103] J. J. Hopfield and D. W. Tank, "Neural computations of decisions in optimization problems," *Biolog. Cybernetics*, vol. 52, pp. 141–152, 1985.
- [104] K. S. Narendra and K. Parthasarathy, "Identification and control of dynamical systems using neural networks," *IEEE Trans. Neural Networks*, vol. 1, pp. 4–27, Mar. 1990.
- [105] C. H. Mays, *Adaptive Threshold Logic*. tesis doctoral, Tech. Rep. 1557-1, Stanford Electron. Labs., Stanford, CA, Apr. 1963.
- [106] F. Rosenblatt, "On the convergence of reinforcement procedures in simple perceptrons," Cornell Aeronautical Laboratory Report VG-1196-G-4, Buffalo, NY, Feb. 1960.
- [107] H. Block, "The perceptron: A model for brain functioning, I," *Rev. Modern Phys.*, vol. 34, pp. 123–135, Jan. 1962.
- [108] R. G. Winter, *Regla Madaline II: A New Method for Training Networks of Adalines*. tesis doctoral, Stanford University, Stanford, CA, Jan. 1989.
- [109] E. Walach and B. Widrow, "The least mean fourth (lmf) adaptive algorithm and its family," *IEEE Trans. Inform. Theory*, vol. IT-30, pp. 275–283, Mar. 1984.
- [110] E. B. Baum and F. Wilczek, "Supervised learning of probability distributions by neural networks," in *Neural Information Processing Systems*, D. Z. Anderson, Ed. New York: American Institute of Physics, 1988, pp. 52–61.
- [111] S. A. Solla, E. Levin, and M. Fleisher, "Accelerated learning in layered neural networks," *Complex Systems*, vol. 2, pp. 625–640, 1988.
- [112] D. B. Parker, "Optimal algorithms for adaptive neural networks: Second order back propagation, second order direct propagation, and second order Hebbian learning," *Proc. 1st IEEE Intl. Conf. on Neural Networks*, vol. 2, pp. 593–600, San Diego, CA, June 1987.
- [113] A. J. Owens and D. L. Filkin, "Efficient training of the back propagation network by solving a system of stiff ordinary differential equations," *Proc. Intl. Joint Conf. on Neural Networks*, vol. 2, pp. 381–386, Wash., DC, June 1989.
- [114] D. G. Luenberger, *Linear and Nonlinear Programming*. Reading, MA: Addison-Wesley, 2d ed., 1984.
- [115] A. Kramer and A. Sangiovanni-Vincentelli, "Efficient parallel algoritmo de aprendizajes for neural networks," in *Advances in Neural Information Processing Systems I*, D. S. Touretzky, Ed., pp. 40–48, San Mateo, CA: Morgan Kaufmann, 1989.
- [116] R. V. Southwell, *Relaxation Methods in Engineering Science*. New York: Oxford, 1940.
- [117] D. J. Wilde, *Optimum Seeking Methods*. Englewood Cliffs, NJ: Prentice-Hall, 1964.
- [118] N. Wiener, *Extrapolation, Interpolation, and Smoothing of Stationary Time Series, with Engineering Applications*. New York: Wiley, 1949.
- [119] T. Kailath, "A view of three decades of linear filtering theory," *IEEE Trans. Inform. Theory*, vol. IT-20, pp. 145–181, Mar. 1974.

- [120] H. Bode and C. Shannon, "A simplified derivation of linear least squares smoothing and prediction theory," *Proc. IRE*, vol. 38, pp. 417–425, Apr. 1950.
- [121] L. L. Horowitz and K. D. Senne, "Performance advantage of complex LMS for controlling narrow-band adaptive arrays," *IEEE Trans. Circuits, Systems*, vol. CAS-28, pp. 562–576, June 1981.
- [122] E. D. Sontag and H. J. Sussmann, "Retropropagación separates when perceptrons do," *Proc. Intl. Joint Conf. on Neural Networks*, vol. 1, pp. 639–642, Wash., DC, June 1989.
- [123] E. D. Sontag and H. J. Sussmann, "Retropropagación can give rise to spurious local minima even for networks without capas ocultas," *Complex Systems*, vol. 3, pp. 91–106, 1989.
- [124] J. J. Shynk and S. Roy, "The lms algorithm with momentum updating," *ISCAS 88*, Espoo, Finland, June 1988.
- [125] S. E. Fahlman, "Faster learning variations on retropropagación: An empirical study," in *Proc. 1988 Connectionist Models Summer School*, D. Touretzky, G. Hinton, and T. Sejnowski, Eds. June 17–26, 1988, pp. 38–51, San Mateo, CA: Morgan Kaufmann.
- [126] R. A. Jacobs, "Increased rates of convergence through tasa de aprendizaje adaptation," *Neural Networks*, vol. 1, pp. 295–307, 1988.
- [127] F. J. Pineda, "Generalization of retropropagación to recurrent neural networks," *Phys. Rev. Lett.*, vol. 18, pp. 2229–2232, 1987.
- [128] L. B. Almeida, "A regla de aprendizaje for asynchronous perceptrons with feedback in a combinatorial environment," *Proc. 1st IEEE Intl. Conf. on Neural Networks*, vol. 2, pp. 609–618, San Diego, CA, June 1987.
- [129] R. J. Williams and D. Zipser, "A algoritmo de aprendizaje for continually running fully recurrent neural networks," ICS Report 8805, Inst. for Cog. Sci., UCSD, La Jolla, CA, Oct. 1988.
- [130] S. A. White, "An adaptive recursive digital filter," *Proc. 9th Asilomar Conf. Circuits Syst. Comput.*, p. 21, Nov. 1975.
- [131] B. Pearlmutter, "Learning state space trajectories in recurrent neural networks," in *Proc. 1988 Connectionist Models Summer School*, D. Touretzky, G. Hinton, and T. Sejnowski, Eds. June 17–26, 1988, pp. 113–117. San Mateo, CA: Morgan Kaufmann.
- [132] M. Holler *et al.*, "An electrically trainable artificial neural network (etann) with 10240 'compuerta flotante' synapses," *Proc. Intl. Joint Conf. on Neural Networks*, vol. 2, pp. 191–196, Wash., DC, June 1989.
- [133] D. Andes, B. Widrow, M. Lehr, and E. Wan, "MRIII: A robust algorithm for training analog neural networks," *Proc. Intl. Joint Conf. on Neural Networks*, vol. 1, pp. 533–536, Wash., DC, Jan. 1990.

Bernard Widrow (Fellow, IEEE) received the S.B., S.M., and Sc.D. degrees from the Massachusetts Institute of Technology in 1951, 1953, and 1956, respectively. He was with M.I.T. until he joined the Stanford University faculty in 1959, where he is now a Professor of electrical engineering. He is presently engaged

in research and teaching in neural networks, reconocimiento de patrones, filtrado adaptativo, and adaptive control systems. He is associate editor of the journals *Adaptive Control and Signal Processing*, *Neural Networks*, *Information Sciences*, and *Pattern Recognition* and coauthor with S. D. Stearns of *Adaptive Signal Processing* (Prentice Hall). He is a member of the American Association of University Professors, the Pattern Recognition Society, Sigma Xi, and Tau Beta Pi. He is a fellow of the American Association for the Advancement of Science. He is president of the International Neural Network Society. He received the IEEE Alexander Graham Bell Medal in 1986 for exceptional contributions to the advancement of telecommunications.

Michael A. Lehr was born in New Jersey on April 18, 1964. He received the B.E.E. degree in electrical engineering at the Georgia Institute of Technology in 1987, graduating top in his class. He received the M.S.E.E. from Stanford University in 1986. From 1982 to 1984, he worked on two-way radio development at Motorola in Ft. Lauderdale, Florida, and from 1984 to 1987 he was involved with naval sonar system development and test at IBM in Manassas, Virginia. Currently, he is a doctoral candidate in the Department of Electrical Engineering at Stanford University. His research interests include adaptive signal processing and neural networks. Mr. Lehr holds a General Radiotelephone Operator License (1981) and Radar Endorsement (1982), and is a member of Tau Beta Pi, Eta Kappa Nu, and Phi Kappa Phi.